



# 机器学习与人工智能 Machine Learning and Artificial Intelligence

## Lecture 2 Regressions

Yingjie Zhang (张颖婕)

Peking University

yingjiezhang@gsm.pku.edu.cn


2021 Fall

# Group Project

- Sign up the team by **Sept 26** (email/msg TAs)
- Proposal: Due on Oct 10, 2021, 11:59pm (one submission per team)
  - Team members (at most 5 students)
  - Project goals (with a real-world business question and available datasets)
  - Models to address the questions (at least one supervised and one unsupervised learning models)
  - Advanced model applications and deeper analyses are encouraged
- In-Class Presentation (Nov 25)
- Final reports (Due Dec 9, 2021)

# Classification vs. Regression

- **Classification**: the goal is to predict a *class label*, which is a choice from a predefined list of possibilities
- **Regression**: the goal is to predict a continuous number, or a *floating-point number (real number)* in programming (math) terms




# Regression

- Given the value of an input  $X$ , the output  $Y$  belongs to the set of real value  $R$ .
- Evaluation: predict output accurately
- Examples:
  - Predict housing price
  - Forecast precipitation

# Regression


**Example:** Dataset with only one feature  $x$  and one scalar output  $y$



**Q:** What is the function that best fits these points?

# k-NN Regression

**Example:** Dataset with only one feature  $x$  and one scalar output  $y$



$k = 1$

- *Train*: store all  $(x, y)$  pairs
- *Predict*: pick the nearest  $x$  in the training data and return its  $y$

$k = 2$  Nearest Neighbor Distance Weighted Regression

- *Train*: store all  $(x, y)$  pairs
- *Predict*: pick the nearest two instances  $x^{(n1)}$  and  $x^{(n2)}$  in training data and return the weighted average of their  $y$  values

# Linear Regression

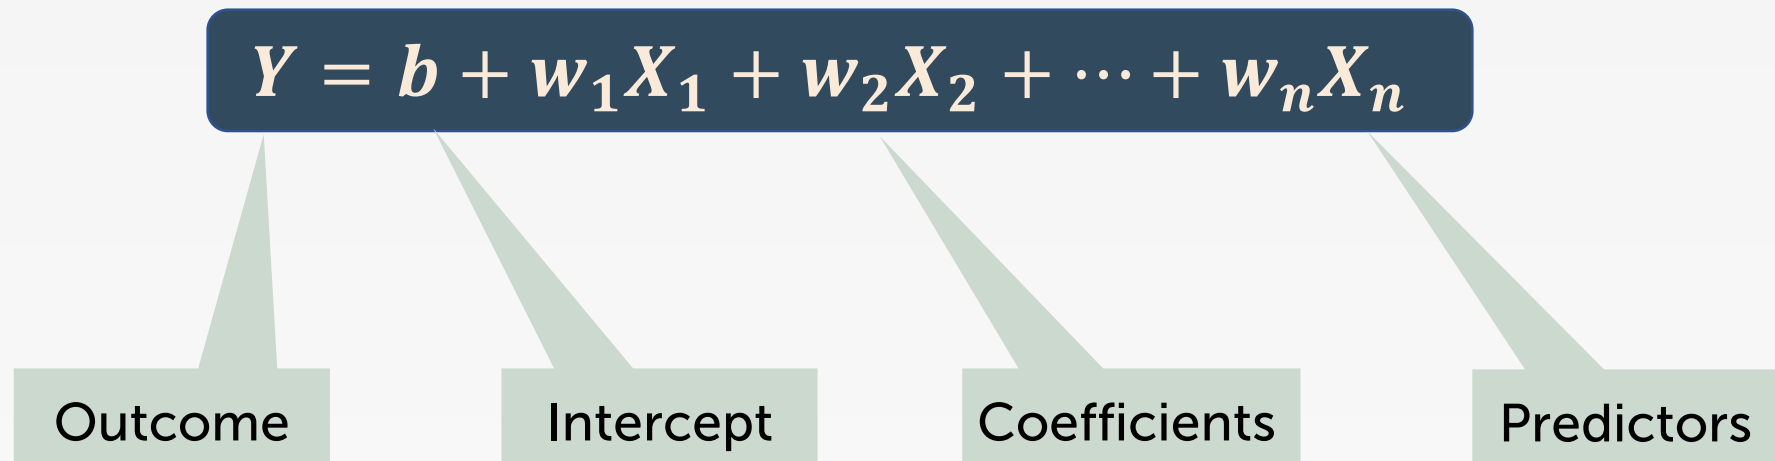
# Agenda

- Definition of Regression
- Linear functions
- Residuals
- Estimations (Optimization)
- Regularization





# Linear Regression

- Linear relationship: outcome (dependent) variable is a linear combination of predictor (independent) variables.





# Linear Model?

$$Y_i = b + wX_i$$





# Linear Model?

$$Y_i = b + wX_i$$





$$Y_i = b + w_1X_i + w_2X_i^2$$





# Linear Model?

$$Y_i = b + wX_i$$



$$Y_i = b + w_1X_i + w_2X_i^2$$



**A linear model means linear in parameters (not X)!**

# Linear Regression?

- $y = \sum_i \omega_i f_i(x)$
- $y = \sum_i \omega_i x^i$
- $y = \sum_i e^{w_i x}$
- $y = \sum_i w_i \sin(i^2 x^7)$

# Results Interpretation

Level-Level

$$y = b + wx + \varepsilon$$

One unit increase of  $x$   
 $\rightarrow w$  unit increase of  $y$

# Results Interpretation

Level-Level

$$y = b + wx + \varepsilon$$

One unit increase of  $x$   
 $\rightarrow w$  unit increase of  $y$

Log-Level

$$\log(y) = b + wx + \varepsilon$$

One unit increase of  $x \rightarrow$   
 $100w\%$  increase of  $y$

Example:  $y$  – income;  $x$  – tenured year;  $w = 0.04$ . One more tenured year increases 4% in income

# Results Interpretation

Level-Level

$$y = b + wx + \varepsilon$$

One unit increase of  $x$   
 $\rightarrow w$  unit increase of  $y$

Log-Level

$$\log(y) = b + wx + \varepsilon$$

One unit increase of  $x \rightarrow$   
 $100w\%$  increase of  $y$

Example:  $y$  – income;  $x$  – tenured year;  $w = 0.04$ . One more tenured year increases 4% in income

Log-Log

$$\log(y) = b + w \log(x) + \varepsilon$$



One percent increase of  $x$   
 $\rightarrow w\%$  increase of  $y$

Example:  $y$  – demand;  $x$  – price;  $w = -0.6$ . 1% increase in price leads to 0.6% decrease in demand





# Linear Model?

Strong relationship




Weak relationship



How to evaluation?

# Residuals



Residuals  $e = \text{observed (y)} - \text{predicted (y)}$

# Train a Linear Model

- Goal: minimize the Error
- Potential ways:
  - Sum (mean) of absolute errors  $|e_1| + |e_2| + |e_3| + \dots$
  - Sum (mean) of squared errors  $e_1^2 + e_2^2 + e_3^2 + \dots$

# Function Approximation

- Objective function: mean squared error (MSE)

$$J(\boldsymbol{\theta}) = \frac{1}{N} \sum_{i=1}^N e_i^2$$

$$= \frac{1}{N} \sum_{i=1}^N (y^{(i)} - \boldsymbol{\theta}^T \mathbf{x}^{(i)})^2$$

$$\mathbf{x}' = [1, x_1, x_2, \dots, x_M]^T$$

$$\boldsymbol{\theta} = [b, w_1, \dots, w_M]^T$$

- Solve the unconstrained optimization problem
  - Closed form
  - Gradient descent
  - Stochastic gradient descent

$$\hat{\boldsymbol{\theta}} = \underset{\boldsymbol{\theta}}{\operatorname{argmin}} J(\boldsymbol{\theta})$$

- Test time:
  - Given a new  $x$ , make a prediction  $\hat{y} = \hat{\boldsymbol{\theta}}^T x$

# Closed-form Solution

## Criteria

Minimize the sum of squared errors

$$\min \sum_n (y_i - \hat{y}_i)^2$$

## Solution

$$y_i = b + w_1 x_i + \varepsilon_i$$

$w_1$ : slope for the estimated regression equation

$$w_1 = \frac{\sum_n (x_i - \bar{x})(y_i - \bar{y})}{\sum_n (x_i - \bar{x})^2}$$

$b$ : intercept for the estimated regression equation

$$b = \bar{y} - w_1 \bar{x}$$

$$\theta = (X^T X)^{-1} X^T y$$

# Gradient Descent



# Linear Regression Solutions

- Closed-form solution:
  - Computational complexity
  - Stability

$$\beta = (X^T X)^{-1} X^T y$$

- Gradient Descent for Linear Regression

# Contour Plots





# Contour Plots


1. Each level curve labeled with value
2. Value label indicates the value of the function for all points lying on that level curve

# Optimization by Random Guessing

## Random guessing:

1. Pick a random  $\theta$
2. Evaluate  $J(\theta)$
3. Repeat steps 1 and 2 many times
4. Return  $\theta$  that gives smallest  $J(\theta)$

$$J(\theta) = J(\theta_1, \theta_2) = (10(\theta_1 - 0.5))^2 + (6(\theta_2 - 0.4))^2$$



# Optimization by Random Guessing


## Random guessing:

1. Pick a random  $\theta$
2. Evaluate  $J(\theta)$
3. Repeat steps 1 and 2 many times
4. Return  $\theta$  that gives smallest  $J(\theta)$

## Linear Regression:


1. Objective function: MSE
2. contour plot: each line labeled with MSE – lower means a better fit
3. **minimum** corresponds to parameters  $(w, b) = (\theta_1, \theta_2)$  that **best fit** some training dataset


$$J(\theta) = J(\theta_1, \theta_2) = (10(\theta_1 - 0.5))^2 + (6(\theta_2 - 0.4))^2$$





# Gradient Descent








# Pros and Cons


- Advantages:
  - Simple and often quite effective on ML tasks
  - Often very scalable
- Drawbacks
  - Might find a local minimum
  - Only applies to smooth function (differentiable)



# Algorithm

## Algorithm 1 Gradient Descent


```
1: procedure GD( $\mathcal{D}$ ,  $\theta^{(0)}$ )  
2:    $\theta \leftarrow \theta^{(0)}$   
3:   while not converged do  
4:      $\theta \leftarrow \theta - \gamma \nabla_{\theta} J(\theta)$   
5:   return  $\theta$ 
```




Convergence Criteria (one example):  $\|\nabla_{\theta} J(\theta)\|_2 \leq \epsilon$




# Gradient Descent



# Learning Rate




# GD for Linear Regression




$t$	$\theta_1$	$\theta_2$	$J$

$$J(\boldsymbol{\theta}) = J(\theta_1, \theta_2) = \frac{1}{N} \sum (y^{(i)} - \boldsymbol{\theta}^T \mathbf{x}^{(i)})^2$$




# GD for Linear Regression




$t$	$\theta_1$	$\theta_2$	$J$
1	0.01	0.02	25.2

$$J(\boldsymbol{\theta}) = J(\theta_1, \theta_2) = \frac{1}{N} \sum (y^{(i)} - \boldsymbol{\theta}^T \mathbf{x}^{(i)})^2$$




# GD for Linear Regression




$t$	$\theta_1$	$\theta_2$	$J$
1	0.01	0.02	25.2
2	0.30	0.12	8.7

$$J(\theta) = J(\theta_1, \theta_2) = \frac{1}{N} \sum (y^{(i)} - \theta^T x^{(i)})^2$$




# GD for Linear Regression




$t$	$\theta_1$	$\theta_2$	$J$
1	0.01	0.02	25.2
2	0.30	0.12	8.7
3	0.51	0.30	1.5

$$J(\theta) = J(\theta_1, \theta_2) = \frac{1}{N} \sum (y^{(i)} - \theta^T x^{(i)})^2$$




# GD for Linear Regression




$t$	$\theta_1$	$\theta_2$	$J$
1	0.01	0.02	25.2
2	0.30	0.12	8.7
3	0.51	0.30	1.5
4	0.59	0.43	0.2

$$J(\theta) = J(\theta_1, \theta_2) = \frac{1}{N} \sum (y^{(i)} - \theta^T x^{(i)})^2$$





# GD for Linear Regression



$t$	$\theta_1$	$\theta_2$	$J$
1	0.01	0.02	25.2
2	0.30	0.12	8.7
3	0.51	0.30	1.5
4	0.59	0.43	0.2

$$J(\theta) = J(\theta_1, \theta_2) = \frac{1}{N} \sum (y^{(i)} - \theta^T x^{(i)})^2$$





# GD for Linear Regression

---

**Algorithm 1** GD for Linear Regression

---

```
1: procedure GDLR( $\mathcal{D}, \theta^{(0)}$ )  
2:    $\theta \leftarrow \theta^{(0)}$                                 ▷ Initialize parameters  
3:   while not converged do  
4:      $\mathbf{g} \leftarrow \sum_{i=1}^N (\theta^T \mathbf{x}^{(i)} - y^{(i)}) \mathbf{x}^{(i)}$     ▷ Compute gradient  
5:      $\theta \leftarrow \theta - \gamma \mathbf{g}$                                 ▷ Update parameters  
6:   return  $\theta$ 
```

---


# Stochastic Gradient Descent (SGD)

## Algorithm 2 Stochastic Gradient Descent (SGD)

```
1: procedure SGD ( $D, \theta^{(0)}$ )  
2:    $\theta \leftarrow \theta^{(0)}$   
3:   while not converged do  
4:     for  $i \sim \text{Uniform}(\{1, 2, 3, \dots, N\})$   
5:        $\theta \leftarrow \theta - \gamma \nabla_{\theta} J^{(i)}(\theta)$   
6:   return  $\theta$ 
```

# Stochastic Gradient Descent

- Just picks a random instance (or sampling) in the training set to compute the gradient



# Mini-Batch SGD


- Gradient Descent:
  - Compute true gradient exactly from all  $N$  examples
- Stochastic Gradient Descent (SGD):
  - Approximate true gradient by the gradient of one randomly chosen example
- Mini-Batch SGD:
  - Approximate true gradient by the average gradient of  $K$  randomly chosen examples

# Logistic Regression

# General View

- Logistic regression is a **classification** algorithm
- It predicts the probability of occurrence of an event by fitting data to a *logit* function
- Outcome: categorical variables

# Logistic vs. Linear



# Logistic Regression

- **Data:** Inputs are continuous vectors of length  $M$ . Outputs are discrete.

$$D = \{\mathbf{x}^{(i)}, y^{(i)}\}_{i=1}^N$$

- **Model:** Logistic function applied to dot product of parameters with input vector.

$$p_{\theta}(y = 1|\mathbf{x}) = \frac{1}{1 + \exp(-\boldsymbol{\theta}^T \mathbf{x})}$$

- **Learning:** finds the parameters that minimize some objective function.

$$\boldsymbol{\theta}^* = \underset{\boldsymbol{\theta}}{\operatorname{argmin}} J(\boldsymbol{\theta})$$

- **Prediction:** Output is the most probable class.

$$\hat{y} = \underset{y \in \{0,1\}}{\operatorname{argmax}} p_{\theta}(y|\mathbf{x})$$



# Additional Information

- Estimates of coefficients ( $\theta$ ) are derived through an iterative process called Maximum Likelihood Estimation (MLE)
- If estimated probability  $>$  Cutoff  $\rightarrow$  Classify as class "1"
- Probability of success (Odds)  $\hat{\pi} = P(y = 1|x) = \frac{e^u}{1+e^u}$
- **Odds-Ratio** for success  $\frac{\hat{\pi}}{1-\hat{\pi}} = e^u$
- Log Odds-Ratio  $\ln\left(\frac{\hat{\pi}}{1-\hat{\pi}}\right) = u = b + w_1X_1 + w_2X_2 + \dots$

# MLE

## Principle of Maximum Likelihood Estimation:

Choose the parameters that maximize the likelihood of the data.

$$\theta^{\text{MLE}} = \underset{\theta}{\operatorname{argmax}} \prod_{i=1}^N p(\mathbf{x}^{(i)} | \theta)$$

Maximum Likelihood Estimate (MLE)

MLE tries to allocate as much probability mass as possible to the things we have observed... ..at the expense of the things we have not observed


# Polynomial Regression

# Polynomial Regression

- Generate new features consisting of all polynomial combinations of the original features
- A linear model
- An application of non-linear transformations




$$X = (x_0, x_1) \longrightarrow X' = (x_0, x_1, x_0x_1, x_0^2, x_1^2)$$
$$Y = w_0x_0 + w_1x_1 + w_{01}x_0x_1 + w_{00}x_0^2 + w_{11}x_1^2$$

# Example I: Polynomial Features




$$\text{True: } y = \tanh(x) + \text{randn}$$

# Example I: Polynomial Features






$$\text{True: } y = \tanh(x) + \text{randn}$$

# Example II: Polynomial Features



$$\text{True: } y = -1.5 \cdot x + \text{randn}$$

# Example II: Polynomial Features




$$\text{True: } y = -1.5 \cdot x + \text{randn}$$




# Overfitting

True:

$$y = -1.5 \cdot x + \text{randn}$$



**Overfitting Definition:** when the model captures the noise in the training data instead of the underlying structure.



# Regularization

# Regularization

- Goal: optimize some combination of fit and simplicity
  - Penalize the magnitude of coefficients of features
  - Minimize the error between predicted and actual examples
- Ridge Regression:
  - L2-norm: adds penalty equivalent to square of the magnitude of coefficients
- Lasso Regression:
  - L1-norm: adds penalty equivalent to absolute value of the magnitude of coefficients

# Ridge Regression

- Recall: in a linear regression with the least square estimation

$$RSS = \sum_{i=1}^n (y_i - (\omega \cdot x_i + b))^2$$

- Ridge regression

$$RSS = \sum_{i=1}^n (y_i - (\omega \cdot x_i + b))^2 + \alpha \sum_{j=1}^p \omega_j^2$$

Shrinkage penalty

# Ridge Regression: $\lambda$

- Ridge regression 
$$RSS = \sum_{i=1}^n (y_i - (\omega \cdot x_i + b))^2 + \alpha \sum_{j=1}^p \omega_j^2$$

$\alpha$ : tuning parameter

- $\alpha = 0$ :
  - A simple linear regression
- $\alpha = \infty$ :
  - Coefficients  $\omega$  will be zero
- As  $\alpha$  increases, the flexibility of the model fit decreases

# LASSO

Least Absolute Shrinkage and Selection Operator Regression

- L1 Regularization

$$RSS = \sum_{i=1}^n (y_i - (\omega \cdot x_i + b))^2 + \alpha \sum_{j=1}^p |\omega_j|$$

- Lasso combines some of the shrinking advantages of ridge with **variable selection**
  - The L1 penalty has the effects of forcing some coefficient estimates to be exactly equal to zero

**Tip:** Techniques such as cross validation are recommended to determine which approach is better on a particular dataset

# Questions?