



机器学习与人工智能 Machine Learning and Artificial Intelligence

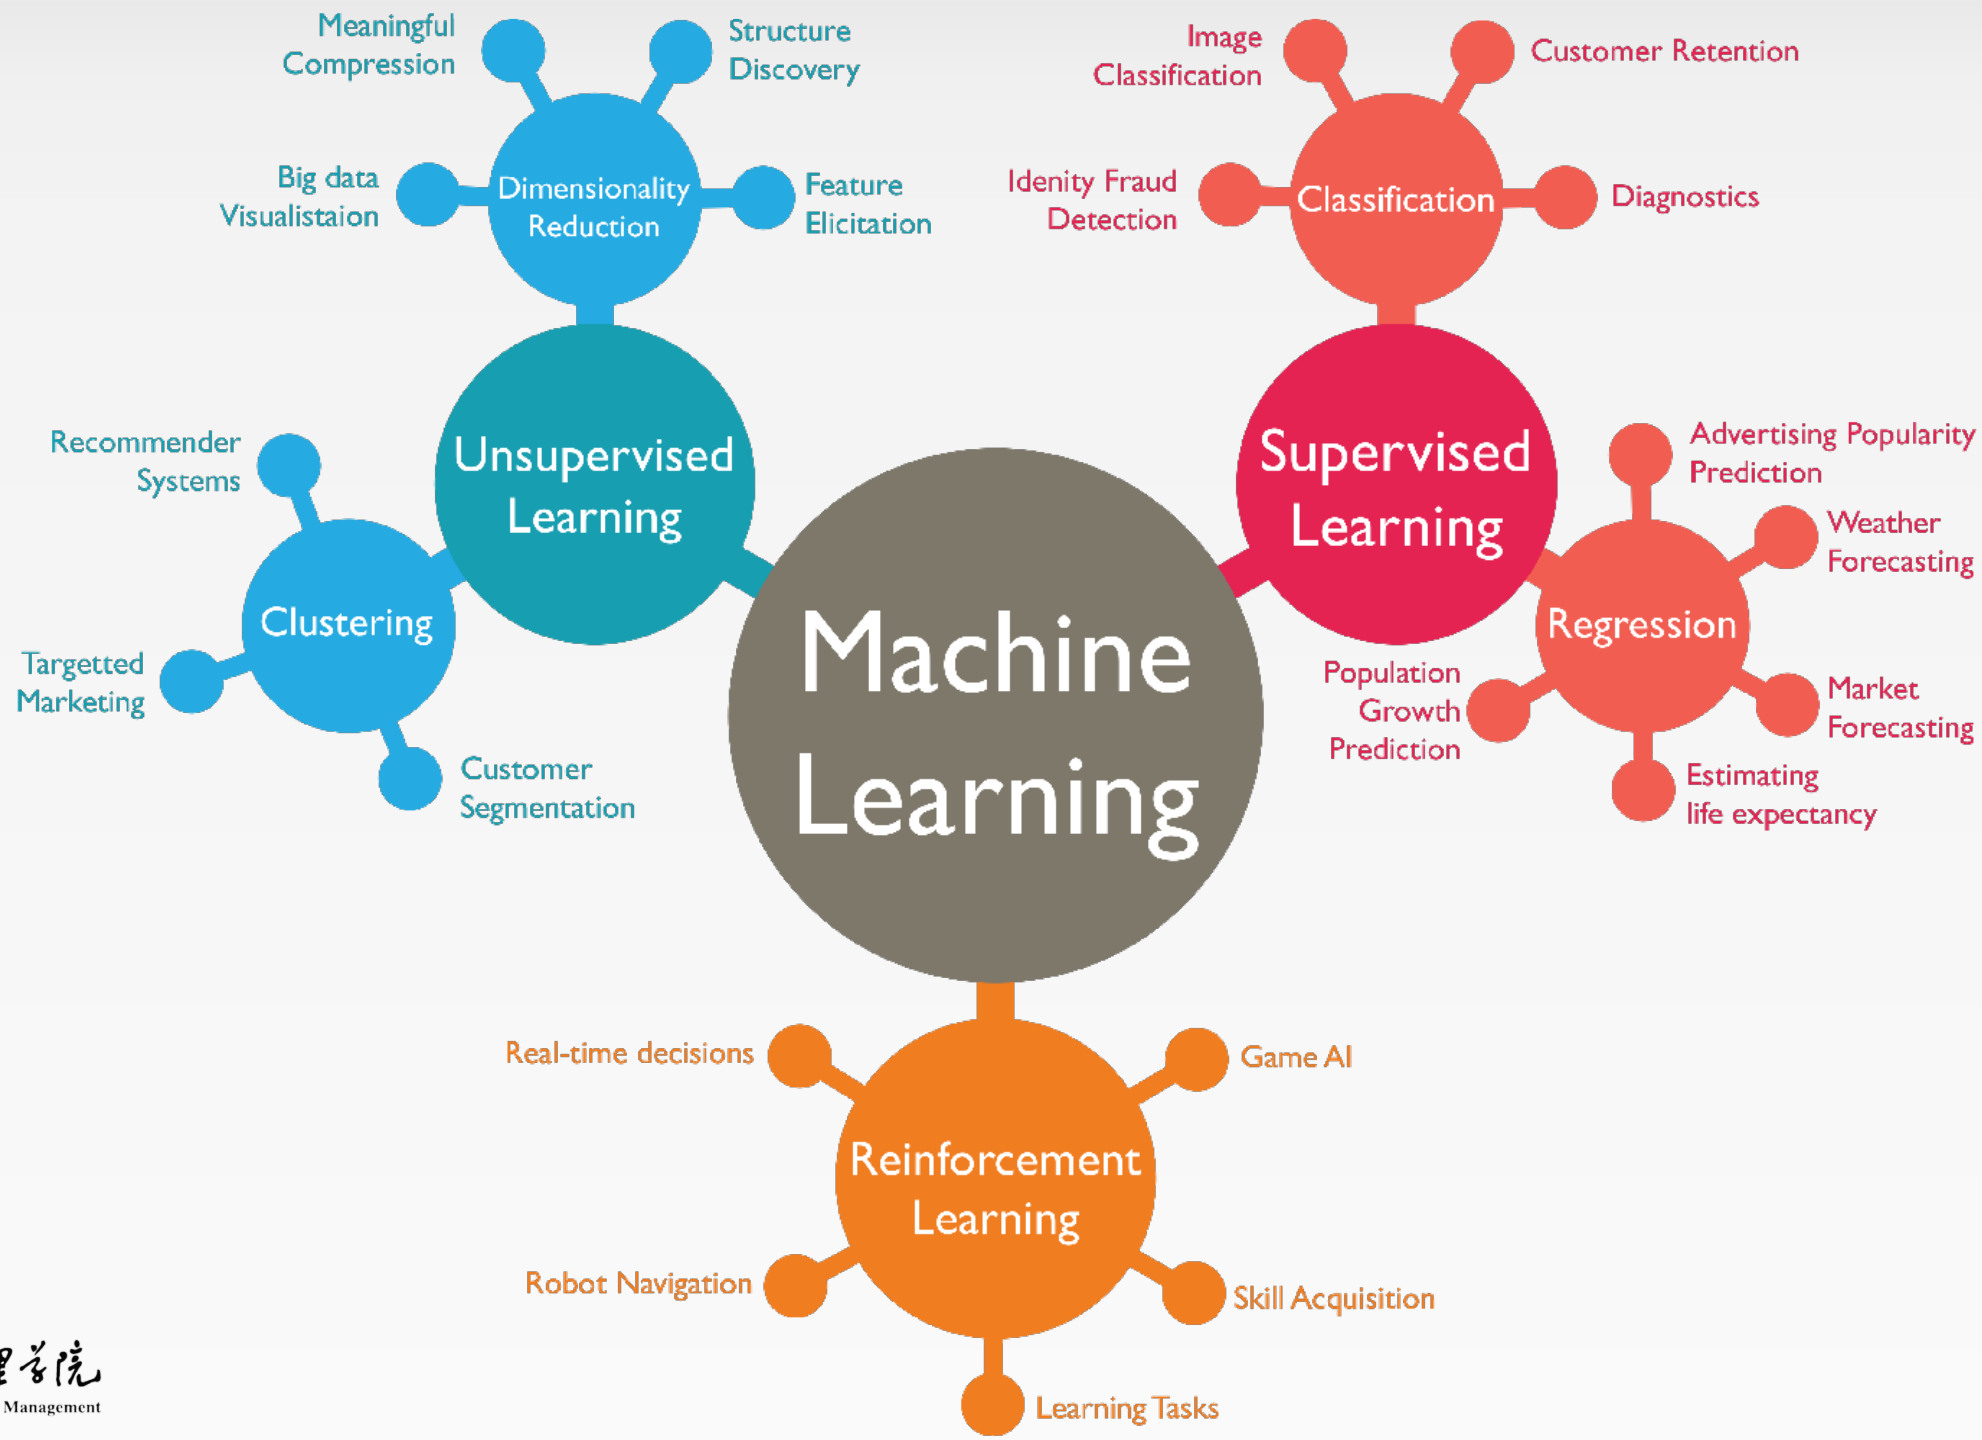
Lecture 6 Ensemble Models,
HMM, Clustering

Yingjie Zhang (张颖婕)

Peking University


yingjiezhang@gsm.pku.edu.cn

2021 Fall



Ensemble Models

Wisdom of the Crowd



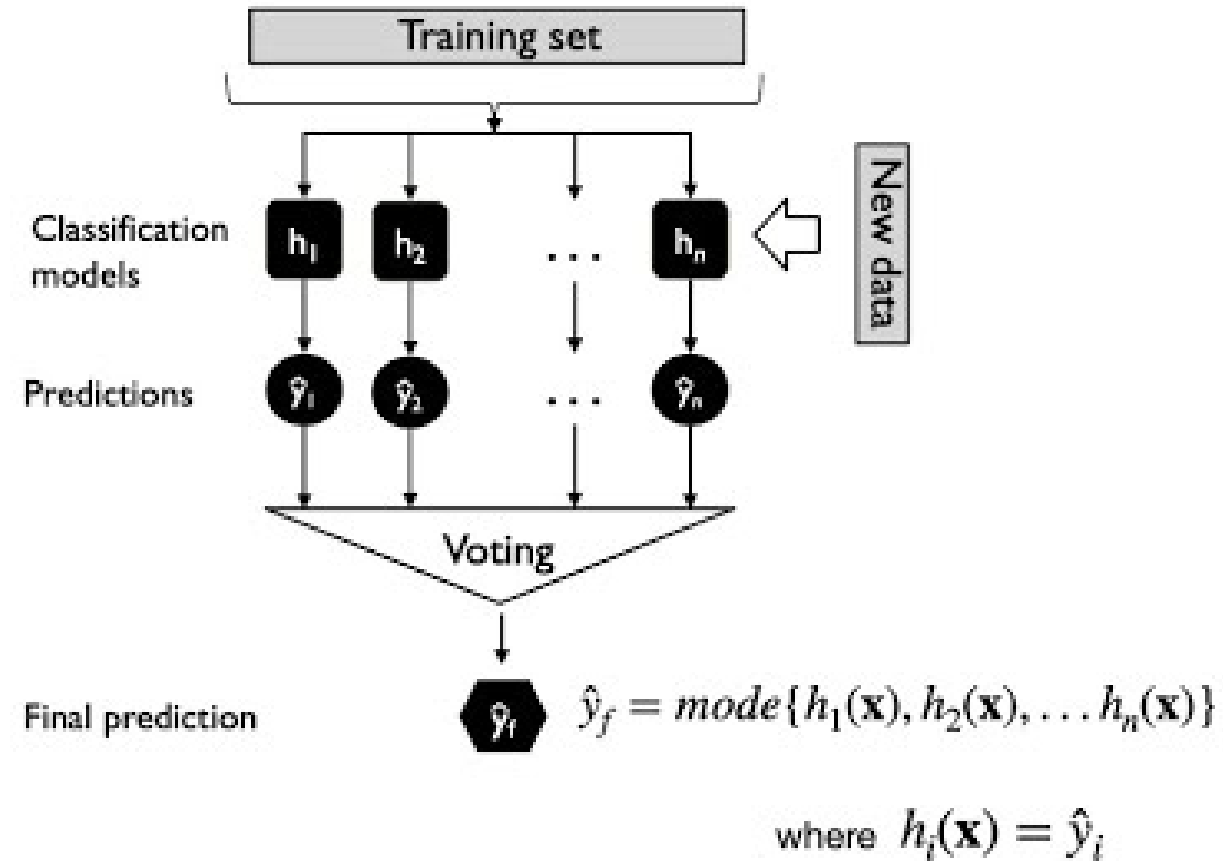
Ensemble Methods

- Why it works:
 - Diversity!
 - Imagine that we have 5 completely independent classifiers; each of them individually is correct 70% of the time
 - Prob(correctly classify a record by a majority vote)
$$= C_{(5,3)}(0.7)^3(0.3)^2 + C_{(5,4)}(0.7)^4(0.3)^1 + C_{(5,5)}(0.7)^5 = 0.837$$
- Downside:
 - Increased complexity, more difficult to interpret
 - Does not always guarantee performance improvements

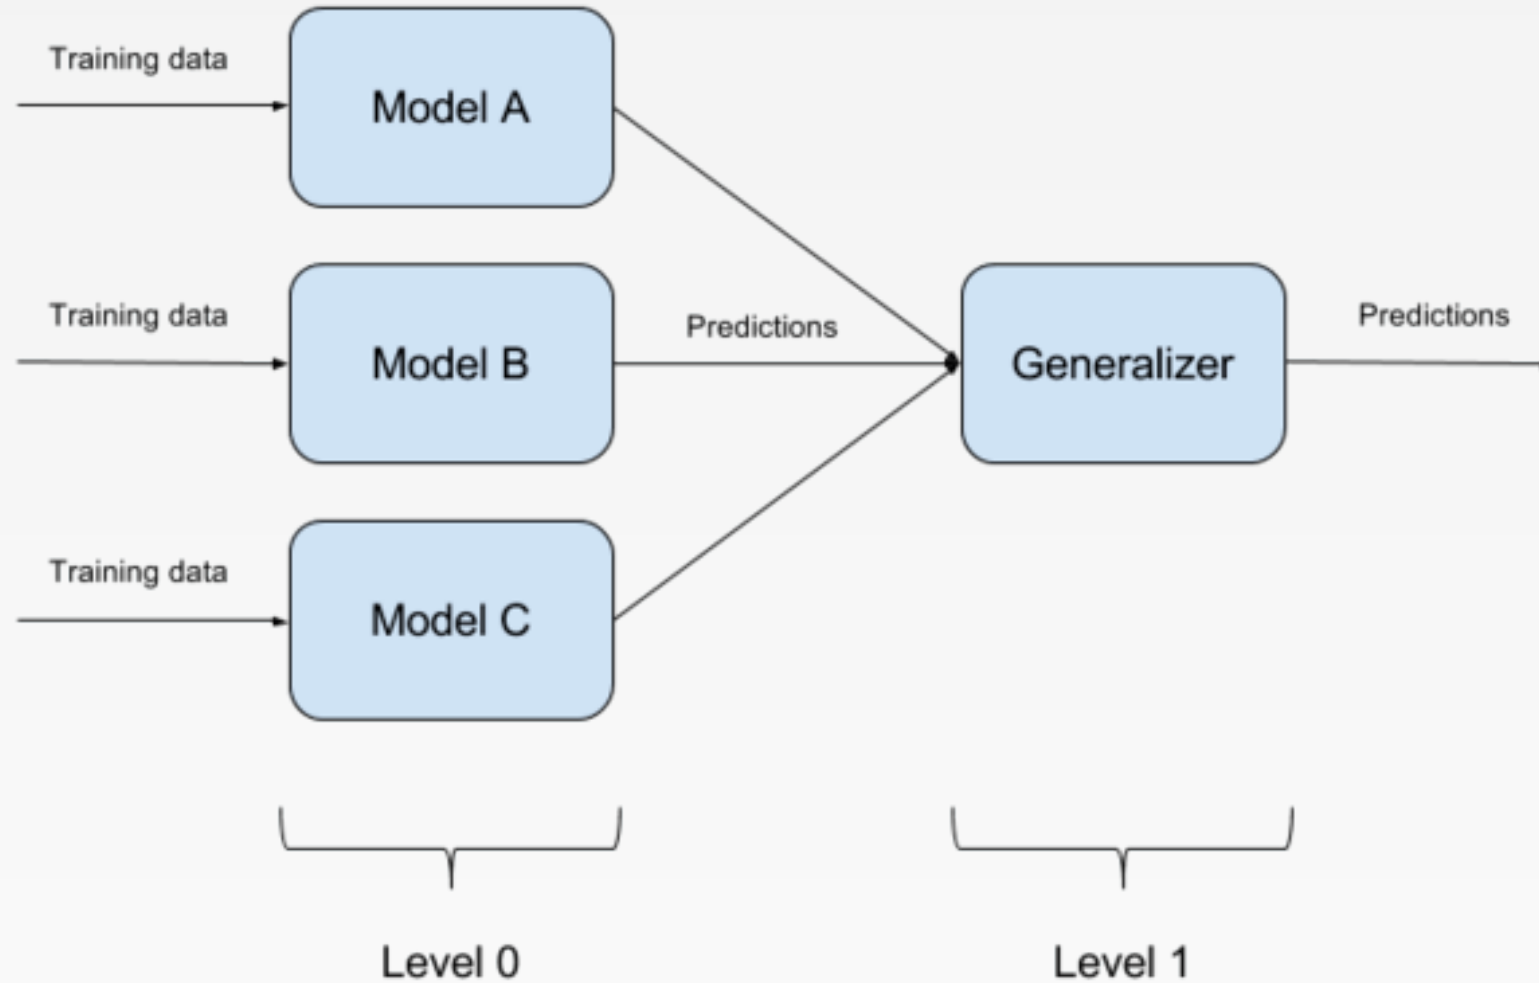
Ensemble Methods

- Voting Classifiers
- Stacking
- Bagging
- Boosting


Voting Classifiers



Stacking




Bagging



Parallel

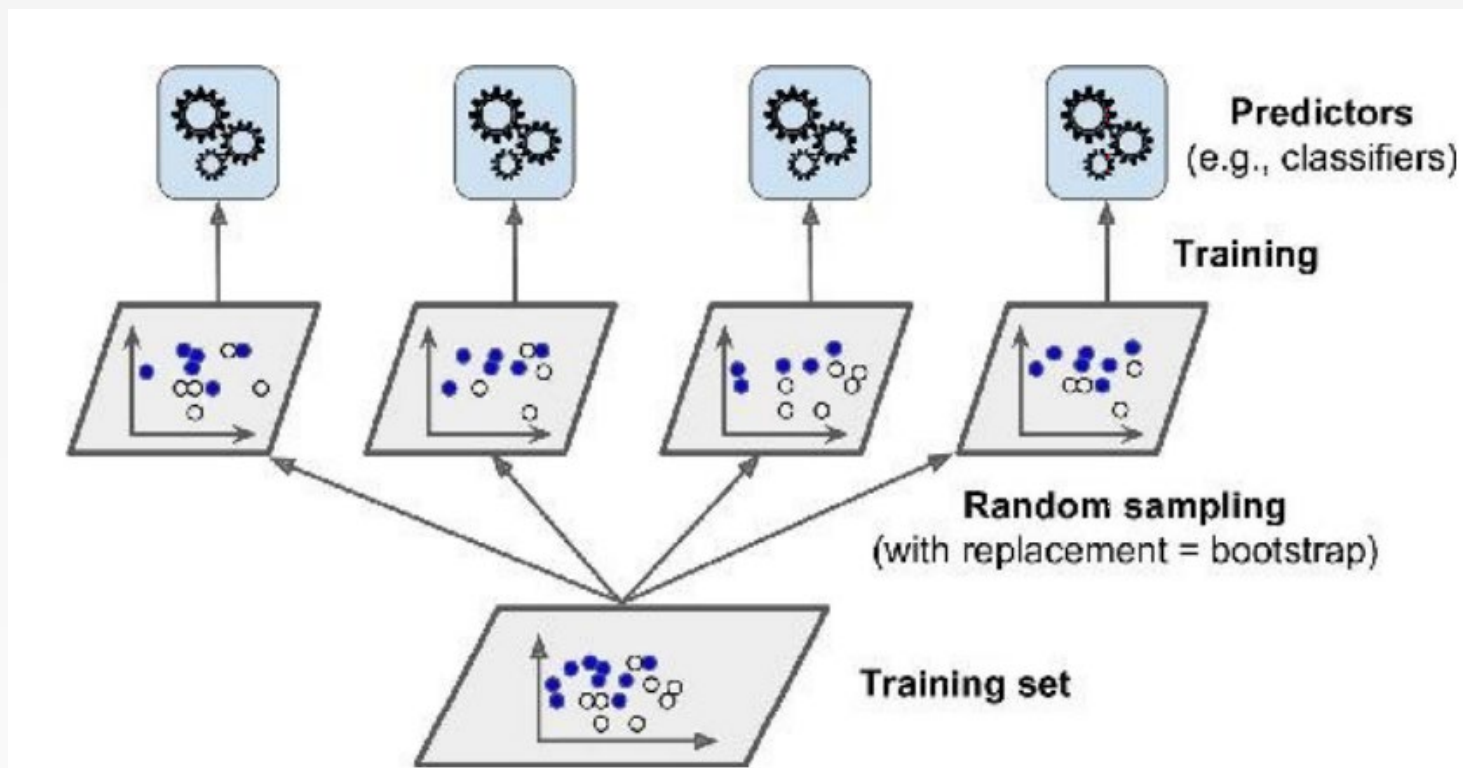
Boosting



Sequential

Bagging: Bootstrap Aggregation

- Ideas:
 - Use the same training algorithm for every predictor, but to train them on different random subsets of the training set



Bagging

- Given
 - Labelled dataset
 - Specific predictive modeling techniques
- Train k models on different training data samples
 - Bootstrap samples: sampled with replacement, typically of the same size as the original training data
- Final prediction is done by combining (i.e., majority vote, averaging) the predictions of k individual models

Overview


- Definition
 - Collection of unpruned trees
 - Rule to combine individual tree decisions
- Purpose
 - Improve prediction accuracy
 - Improve efficiency
- Principle
 - Encouraging diversity among the tree
- Solution: randomness
 - Bagging
 - Random decision trees

Details


- Build many “random” trees
- Randomness: using only a random sample of m attributes to calculate each split
- For each tree:
 - Choose a different training sample
 - For each node, choose m random attributes and find the best split
 - Trees are often fully grown (not pruned)
- Predication: majority vote among all the trees

Boosting


- AdaBoost



AdaBoost: Toy Example




AdaBoost: Toy Example




$$\begin{aligned}\varepsilon_1 &= 0.3 \\ \alpha_1 &= 0.42\end{aligned}$$

AdaBoost: Toy Example




$$\begin{aligned}\varepsilon_1 &= 0.3 \\ \alpha_1 &= 0.42\end{aligned}$$

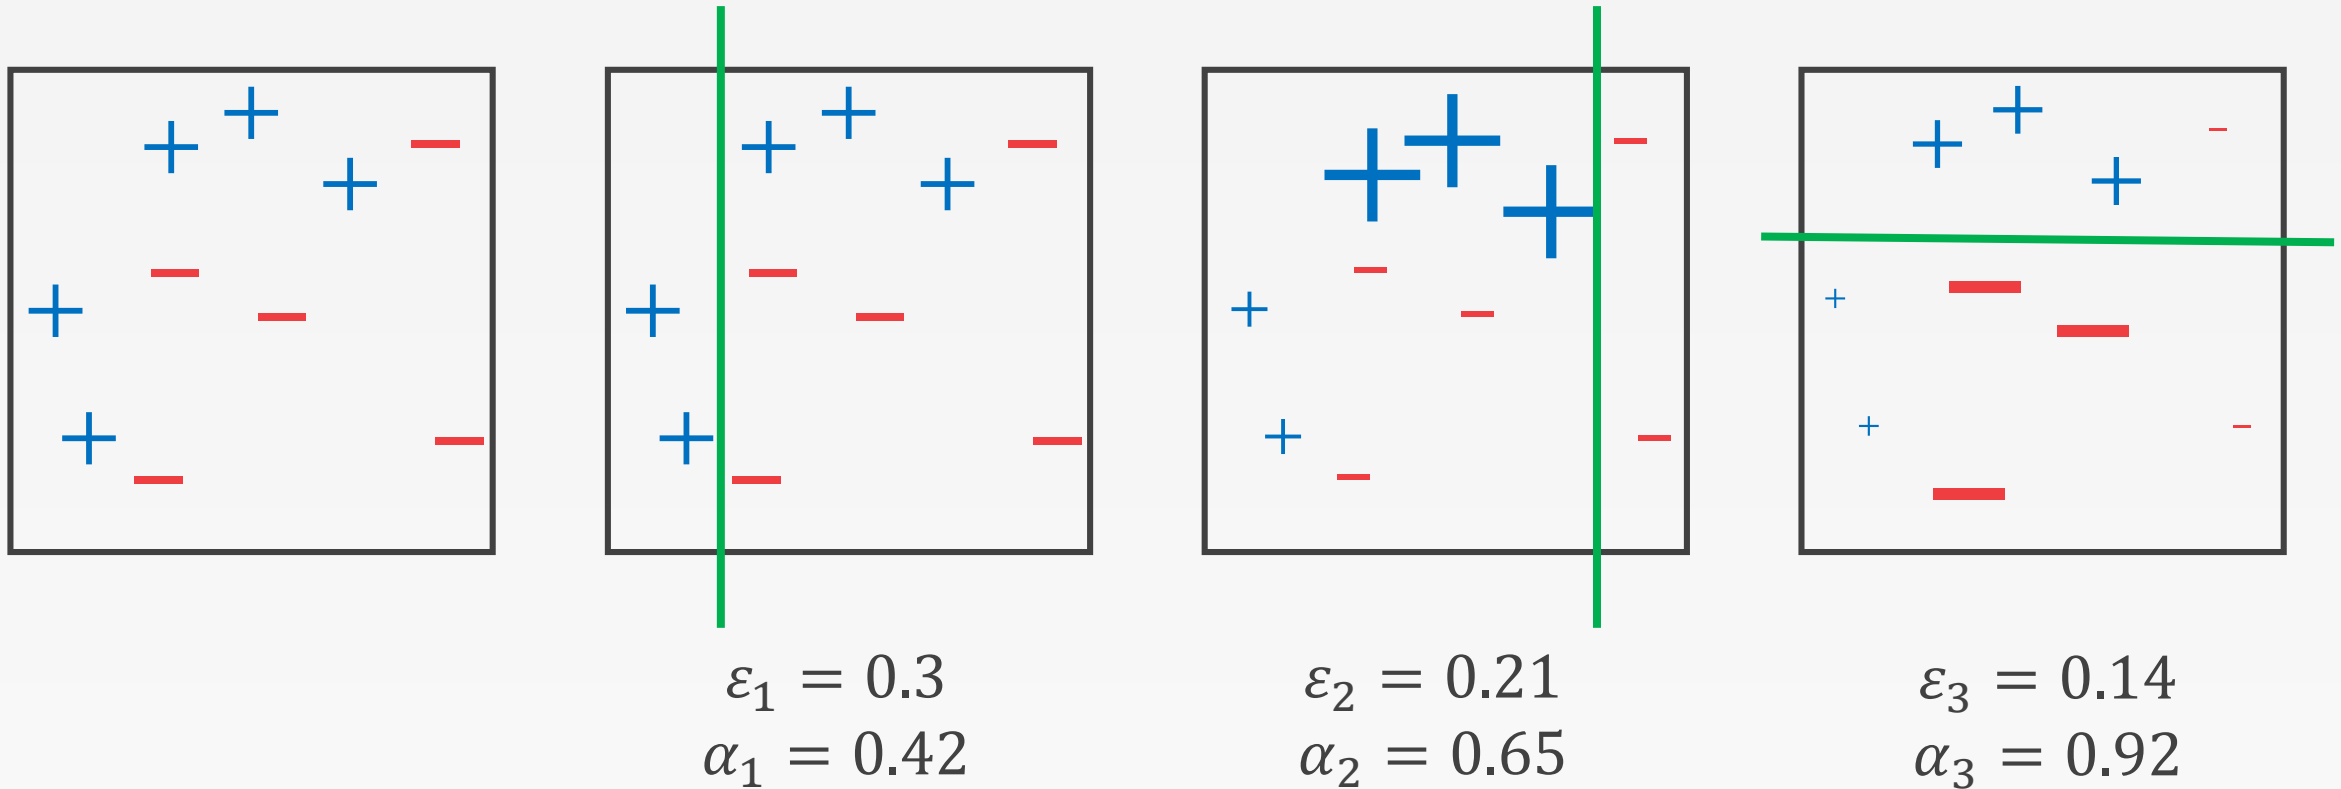
AdaBoost: Toy Example




AdaBoost: Toy Example



AdaBoost: Toy Example



AdaBoost: Toy Example



AdaBoost Algorithm

Given: $(x_1, y_1), \dots, (x_m, y_m)$ where $x_i \in X, y_i \in Y = \{-1, +1\}$

Initialize $D_1(i) = 1/m$

For $t=1, \dots, T$:

Train weak learner using distribution D_t

Get weak hypothesis $h_t: X \rightarrow \{-1, +1\}$ with error $\varepsilon_t = \Pr_{i \sim D_t} [h_t(x_i) \neq y_i]$

Choose $\alpha_t = \frac{1}{2} \ln \left(\frac{1-\varepsilon_t}{\varepsilon_t} \right)$

Update:

$$D_{t+1}(i) = \frac{D_t(i)}{Z_t} \times \begin{cases} e^{-\alpha_t} & \text{if } h_t(x_i) = y_i \\ e^{\alpha_t} & \text{if } h_t(x_i) \neq y_i \end{cases} = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$$

where Z_t is a normalization factor (chosen so that D_{t+1} will be a distribution)

Output the final hypothesis: $H(x) = \text{sign}(\sum_{t=1}^T \alpha_t h_t(x))$


Hidden Markov Models



HMM

- An HMM provides a joint distribution with an assumption of dependence between adjacent states

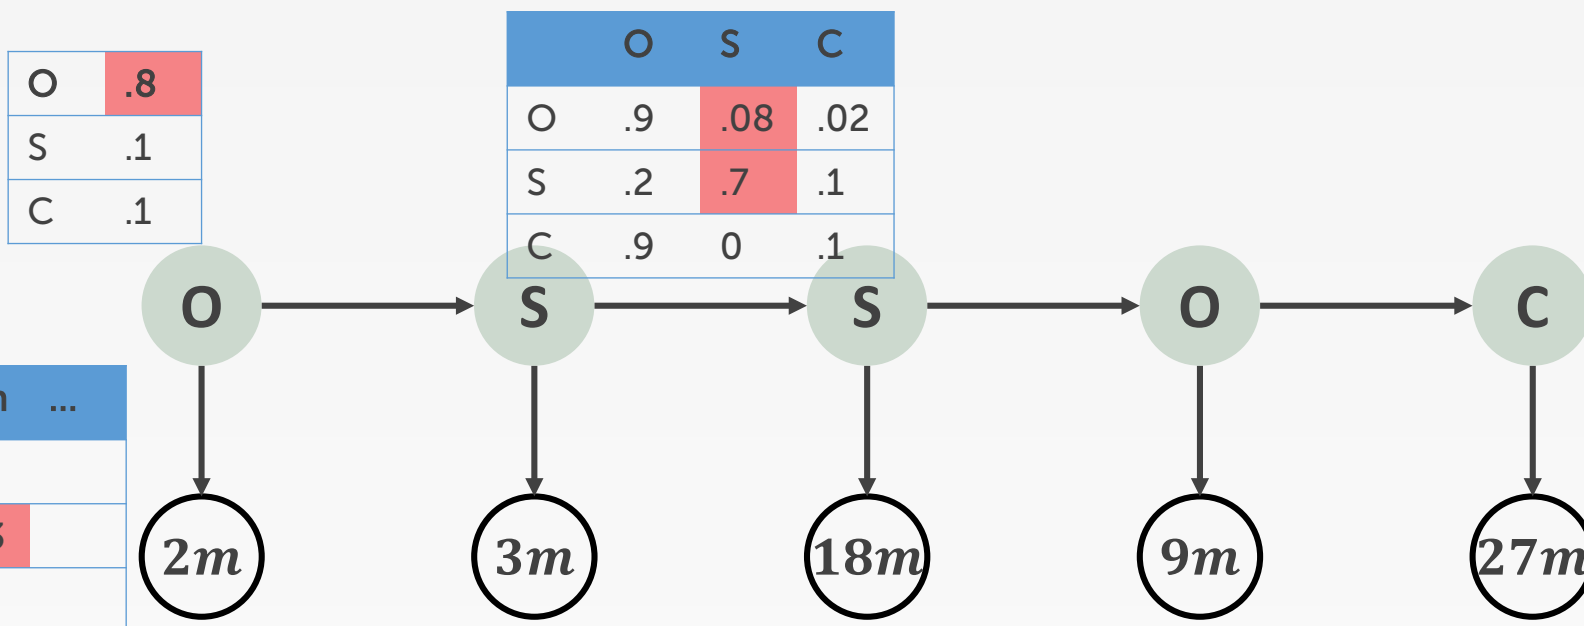
$$p(O, S, S, O, C, 2m, 3m, 18m, 9m, 27m) = (.8 * .2 * .08 * .03 * .7 * \dots)$$




HMM

- An HMM provides a joint distribution with an assumption of dependence between adjacent states

$$p(O, S, S, O, C, 2m, 3m, 18m, 9m, 27m) = (.8 * .2 * .08 * .03 * .7 * \dots)$$




HMM




Naïve Bayes: $P(\mathbf{X}, \mathbf{Y}) = \prod_{t=1}^T P(X_t|Y_t)p(Y_t)$

HMM




Naïve Bayes: $P(\mathbf{X}, \mathbf{Y}) = \prod_{t=1}^T P(X_t|Y_t)p(Y_t)$




HMM: $P(\mathbf{X}, \mathbf{Y}|Y_0) = \prod_{t=1}^T P(X_t|Y_t)p(Y_t|Y_{t-1})$

HMM



Naïve Bayes: $P(\mathbf{X}, \mathbf{Y}) = \prod_{t=1}^T P(X_t|Y_t)p(Y_t)$



HMM: $P(\mathbf{X}, \mathbf{Y}|Y_0) = \prod_{t=1}^T P(X_t|Y_t)p(Y_t|Y_{t-1})$

Supervised Learning for HMM

- HMM Parameters:
 - Emission matrix, A , where $P(X_t = k | Y_t = j) = A_{jk}, \forall t, k$
 - Transition matrix, B , where $P(Y_t = k | Y_{t-1} = j) = B_{jk}, \forall t, k$
- Assumption: $y_0 = START$
- Generative Story:
 - $Y_t \sim \text{Multinomial}(B_{Y_{t-1}}), \forall t$
 - $X_t \sim \text{Multinomial}(A_{Y_t}), \forall t$
- Joint Distribution:
 - $p(\mathbf{X}, \mathbf{Y} | y_0) = \prod_{t=1}^T p(x_t | y_t) p(y_t | y_{t-1}) = \prod_{t=1}^T A_{y_t, x_t} B_{y_{t-1}, y_t}$


Unsupervised Learning for HMMs

- We don't observe any y 's
- This unsupervised learning setting can be achieved by finding parameters that maximize the marginal likelihood
- We optimize using the Expectation-Maximization (EM) algorithm
 - Marginal probability: $p_{\theta}(x) = \sum_{y \in \mathbb{Y}} p_{\theta}(x, y)$
 - $l(\theta) = \log \prod_{i=1}^N p_{\theta}(x^{(i)}) = \sum_{i=1}^N \log \sum_{y \in \mathbb{Y}} p_{\theta}(x^{(i)}, y)$

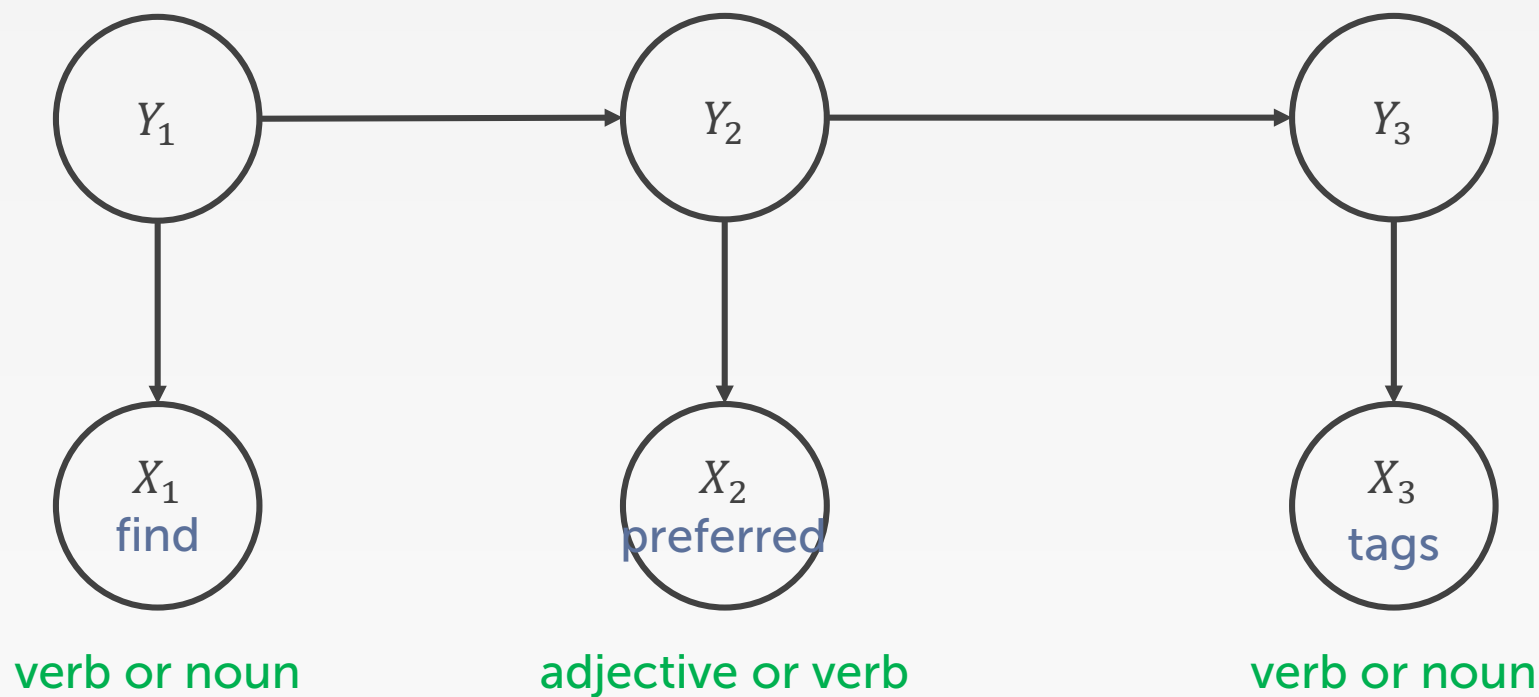
Inference for HMMs

- **Evaluation**: Compute the probability of a given sequence of observations
- **Viterbi Decoding**: Find the most-likely sequence of hidden states, given a sequence of observations
- **Learning**: find the optimal parameters to maximize the probability of the sequence of observations


Part-of-Speech (POS) Tagging




Forward-Backward Algorithm




Forward-Backward Algorithm



Forward-Backward Algorithm



Forward-Backward Algorithm: Finds Marginals



Forward-Backward Algorithm

- Define $\alpha_t(k) \triangleq p(x_1, \dots, x_t, y_t = k)$, $\beta_t(k) \triangleq p(x_{t+1}, \dots, x_T | y_t = k)$
- Assume $y_0 = START$, $y_{T+1} = END$
- 1. Initialize $\alpha_0(START) = \beta_{T+1}(END) = 1$, $\alpha_0(k) = 0, \forall k \neq START$, $\beta_{T+1}(k) = 0, \forall k \neq END$
- 2. Forward algorithm:

for $t = 1, \dots, T$:
 for $k = 1, \dots, K$:

$$\alpha_t(k) = p(x_t | y_t = k) \sum_{j=1}^K \alpha_{t-1}(j) p(y_t = k | y_{t-1} = j)$$
- 3. Backward algorithm:

for $t = T, \dots, 1$:
 for $k = 1, \dots, K$:

$$\beta_t(k) = \sum_{j=1}^K p(x_{t+1} | y_{t+1} = j) \beta_{t+1}(j) p(y_{t+1} = j | y_t = k)$$
- 4. Evaluation: $p(\vec{x}) = \alpha_{T+1}(END)$
- 5. Marginal: $p(y_t = k | \vec{x}) = \frac{\alpha_t(k) \beta_t(k)}{p(\vec{x})}$

Viterbi Algorithm (Decoding)

- Define $\omega_t(k) \triangleq \max_{y_1, \dots, y_{t-1}} p(x_1, \dots, x_t, y_1, \dots, y_t = k)$,

$$b_t(k) \triangleq \operatorname{argmax}_{y_1, \dots, y_{t-1}} p(x_1, \dots, x_t, y_1, \dots, y_t = k)$$

- Assume $y_0 = START$

1. Initialize $\omega_0(START) = 1$, $\omega_0(k) = 0, \forall k \neq START$

2. For $t = 1, \dots, T$:

for $k = 1, \dots, K$:

$$\omega_t(k) = \max_{j \in \{1, \dots, K\}} p(x_t | y_t = k) \omega_{t-1}(j) p(y_t = k | y_{t-1} = j)$$


$$b_t(k) = \operatorname{argmax}_{j \in \{1, \dots, K\}} p(x_t | y_t = k) \omega_{t-1}(j) p(y_t = k | y_{t-1} = j)$$

3. Compute most probable assignment


$$\widehat{y}_T = b_{T+1}(END)$$

$$\text{for } t = T - 1, \dots, 1: \widehat{y}_t = b_{t+1}(\widehat{y}_{t+1})$$


Example: Viterbi Algorithm




Example: Viterbi Algorithm




Example: Viterbi Algorithm



Example: Viterbi Algorithm



Example: Viterbi Algorithm



Unsupervised Learning

Learning Paradigms

Paradigm	Data
Supervised	$\mathcal{D} = \{\mathbf{x}^{(i)}, y^{(i)}\}_{i=1}^N$ $\mathbf{x} \sim p^*(\cdot)$ and $y = c^*(\cdot)$
\hookrightarrow Regression	$y^{(i)} \in \mathbb{R}$
\hookrightarrow Classification	$y^{(i)} \in \{1, \dots, K\}$
\hookrightarrow Binary classification	$y^{(i)} \in \{+1, -1\}$
Unsupervised	$\mathcal{D} = \{\mathbf{x}^{(i)}\}_{i=1}^N$ $\mathbf{x} \sim p^*(\cdot)$
\hookrightarrow Clustering	Predict $\{z^{(i)}\}_{i=1}^N$ where $z^{(i)} \in \{1, \dots, K\}$
\hookrightarrow Dimensionality Reduction	Convert each $\mathbf{x}^{(i)} \in \mathbb{R}^M$ to $\mathbf{u}^{(i)} \in \mathbb{R}^K$ with $K \ll M$
Semi-supervised	$\mathcal{D} = \{\mathbf{x}^{(i)}, y^{(i)}\}_{i=1}^{N_1} \cup \{\mathbf{x}^{(j)}\}_{j=1}^{N_2}$
Reinforcement Learning	$\mathcal{D} = \{(s^{(1)}, a^{(1)}, r^{(1)}), (s^{(2)}, a^{(2)}, r^{(2)}), \dots\}$

Goals

- To discover interesting things from the data:
 - Is there an informative way to visualize the data?
 - Can we discover subgroups among the variables?
- Models:
 - Clustering
 - K-means
 - DBSCAN
 - Hierarchical Clustering

Clustering

Clustering

- Partition **unlabeled** data into groups (clusters)
- Points within a cluster should be “similar”
- Points in different clusters should be “different”

Applications

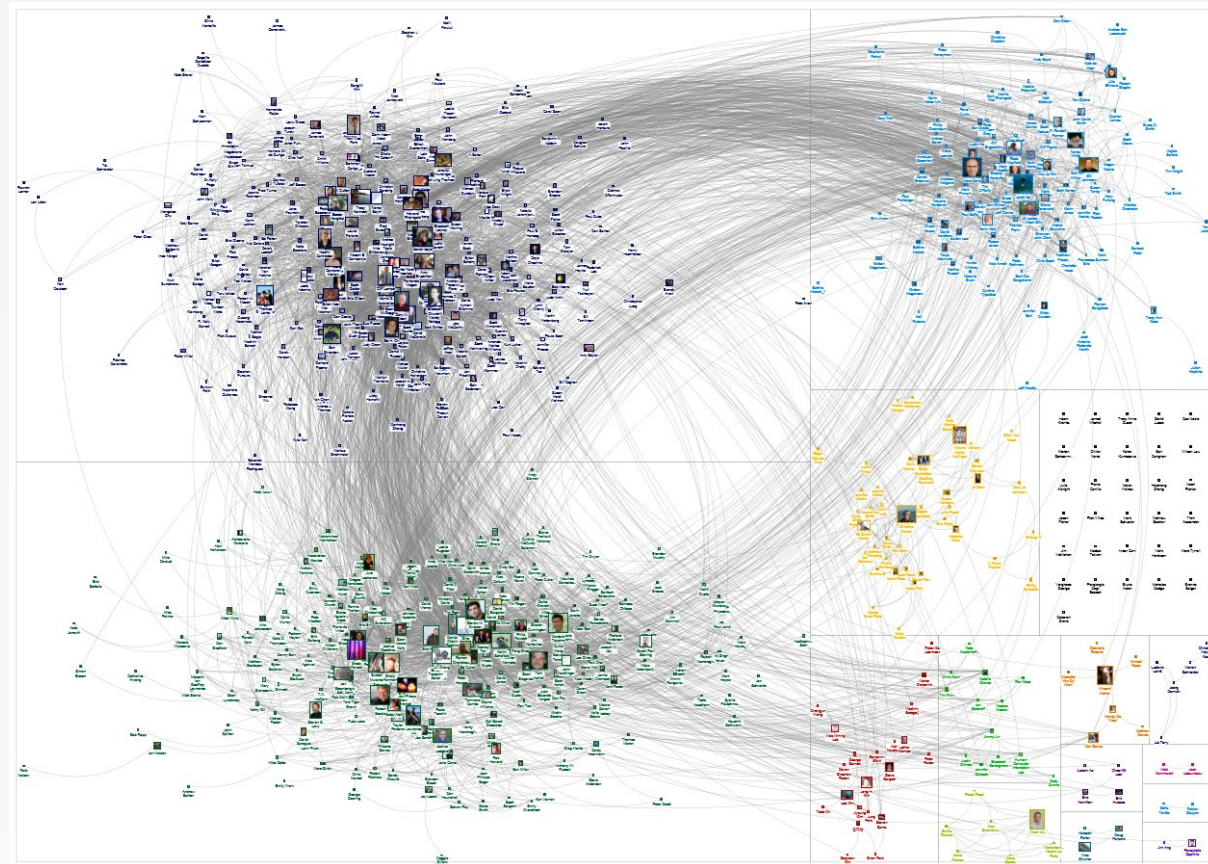
精准货品推荐能力
根据粉丝画像进行商品推荐、实时销量榜单

专享货 全网榜 快手榜 抖音榜
机构专享货 商品热销总榜 快手热销榜 抖音热销榜

1 MAXFACTOR 蜜丝佛陀 魅惑
2小时销量 2781 件
锁佣时间 8.17-9.17
¥189 289 库存 10.5 万
券 100 剩余 400 张
赚 ¥25.87 佣金率 50%

2 Shiseido 资生堂 4色唇膏盘
2小时销量 2781 件
锁佣时间 8.7-9.7
¥189 289 库存 10.5 万
券 100 剩余 400 张
赚 ¥25.56 佣金率 50%

3 2小时销量 2781 件
锁佣时间 8.7-9.7
¥189 289 库存 10.5 万
券 100 剩余 400 张
赚 ¥25.56 佣金率 50%



K-Means

Overview


- K-means (MacQueen, 1967)
- Each cluster has a cluster center, called centroid
- K is specified by the user

K-means Algorithm

- Given K and unlabeled feature vectors $D = \{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(N)}\}$
- Initialize cluster center $\mathbf{c} = \{\mathbf{c}^{(1)}, \dots, \mathbf{c}^{(K)}\}$ and cluster assignments $\mathbf{z} = \{z^{(1)}, z^{(2)}, \dots, z^{(N)}\}$
- Repeat until convergence:
 - For j in $\{1, \dots, K\}$
 $\mathbf{c}^{(j)}$ is the mean of all points assigned to cluster j
 - for i in $\{1, \dots, N\}$
 $z^{(i)}$ is the index j of cluster center nearest to $\mathbf{x}^{(i)}$


Illustrative Example

Given a set of data points




Illustrative Example

Select initial centers at random ($k=3$)




Illustrative Example

Assign each point to its nearest center




Illustrative Example

Recompute optimal centers given a fixed clustering




Illustrative Example

Assign each point to its nearest center




Illustrative Example

Recompute optimal centers given a fixed clustering




Illustrative Example

Assign each point to its nearest center



Illustrative Example

Recompute optimal centers given a fixed clustering



Measure the Distance

- Similarity measure (distance measure)
 - Euclidean distance $d(x, y) = \sqrt{(x - y)^2} = \sqrt{\sum_{i=1}^d (x_i - y_i)^2}$
 - Manhattan distance $d(x, y) = |x - y| = \sum_{i=1}^d |x_i - y_i|$

Stopping Criterion

- no (or minimum) re-assignments of data points to different clusters, *or*
- no (or minimum) change of centroids, or
- minimum decrease in the **sum of squared error(SSE)**,


How to choose k ?

Elbow method:

- run k-means clustering on the dataset for a range of values of k


- for each value of k calculate the sum of squared errors (SSE)



- If the line chart looks like an arm, then the "elbow" on the arm is the value of k that is the best



Pros and Cons

- Strengths:
 - Simple: each to understand and to implement
 - Efficient
- Weakness:
 - The algorithm is sensitive to outliers
 - it terminates at a **local optimum** if SSE is used. The global optimum is hard to find due to complexity
 - Might be sensitive to initial seeds
 - Only simple cluster shapes





DBSCAN

Density-Based Spatial Clustering of Applications with Noise


Density-based Clustering

- Basic Idea:
 - Clusters are dense regions in the data space, separated by regions of lower object density
 - A cluster is defined as a maximal set of density-connected points

Density Definition

- ε -Neighborhood – Objects within a radius of ε from an object

$$N_\varepsilon(p): \{q | d(p, q) \leq \varepsilon\}$$
- “High density” -- ε -Neighborhood of an object contains at least ***MinPts*** of objects

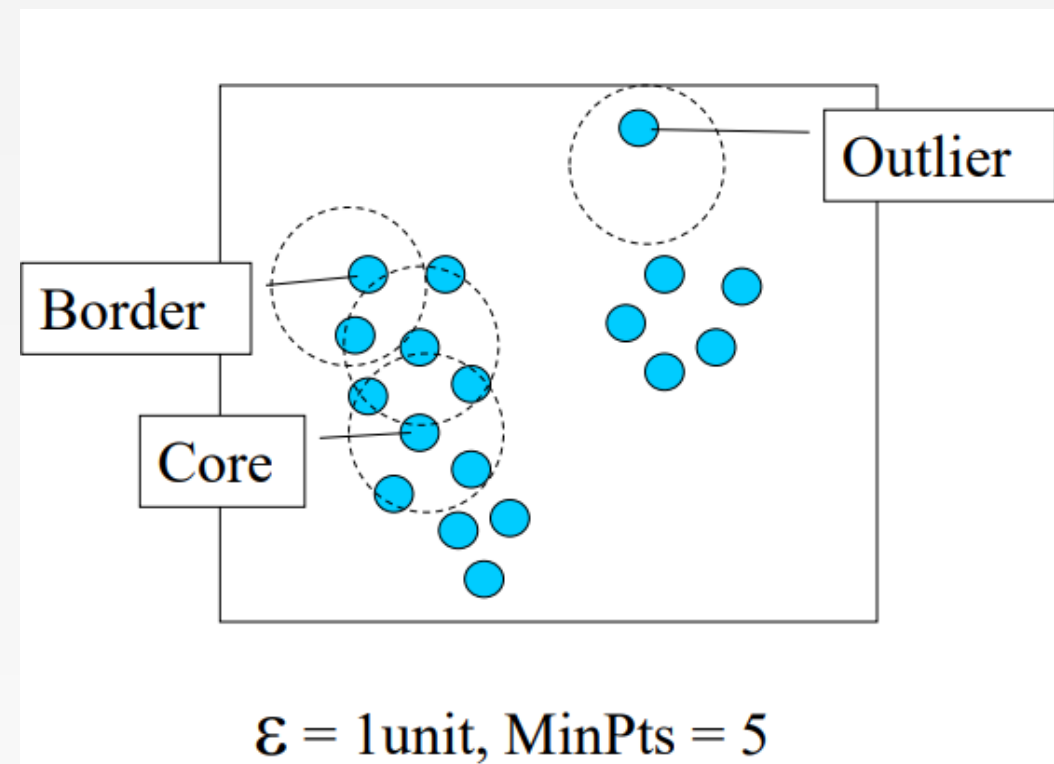


Density of p is “high” ($MinPts = 4$)

Density of q is “low” ($MinPts = 3$)


Core, Border, Outlier

- Given ε and $MinPts$, categorize the objects into three exclusive groups:
 - Core point**: has more than $MinPts$ points within ε (these are points that are at the interior of a cluster)
 - Border point**: has fewer than $MinPts$ within ε , but is the neighborhood of a core point
 - Noise point**: any point that is neither a core nor a border point



Density-reachability


- An object q is directly density-reachable from object p if p is a core object and q is in p 's ε -neighborhood.



MinPts=4

q is directly density-reachable from p
 p is not directly density-reachable from q
Density-reachability is asymmetric

Density-reachability




A point p is directly density-reachable from p_2
 p_2 is directly density-reachable from p_1
 p_1 is directly density-reachable from q
 $p \leftarrow p_2 \leftarrow p_1 \leftarrow q$ form a chain

DBSCAN Algorithm

```
for each  $o \in D$  do
  if  $o$  is not yet classified then
    if  $|o's \ \varepsilon\text{-neighborhood}| < MinPts$ 
      assign  $o$  to NOISE
    else
      collect all objects density-reachable from  $o$ 
      and assign them to a new cluster
```

Pros and Cons

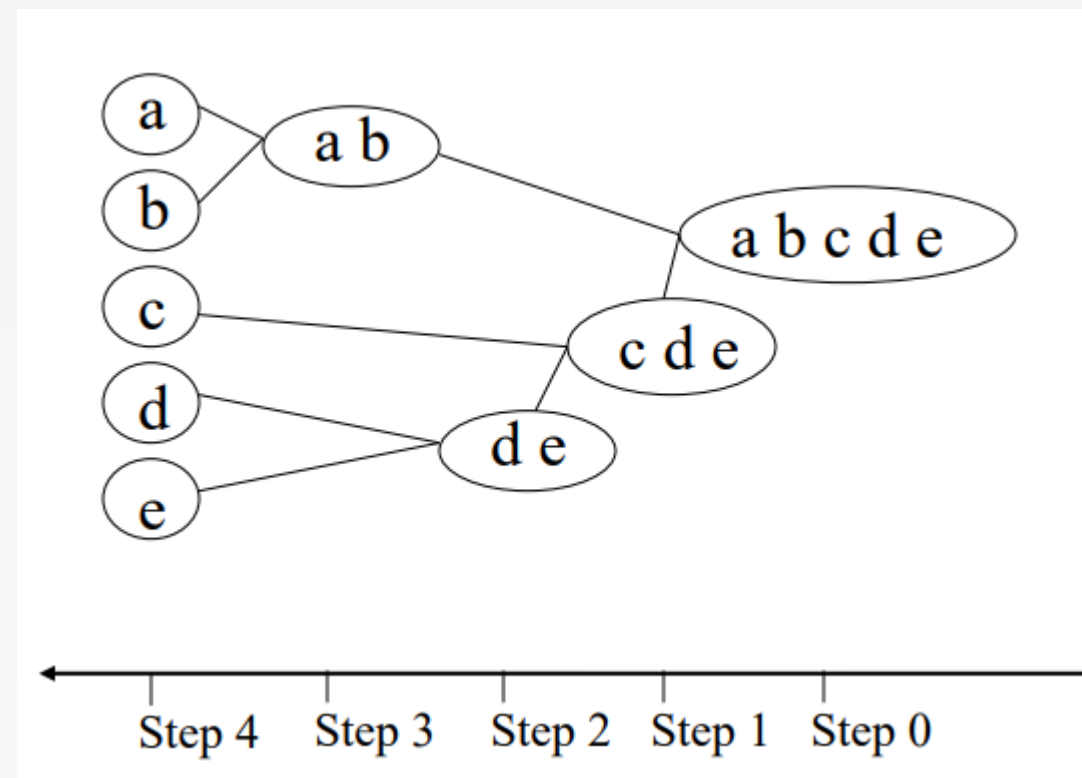
- Can learn arbitrary cluster shapes (resistant to noise)
- Can detect outliers
- Needs two parameters to adjust



Hierarchical Clustering

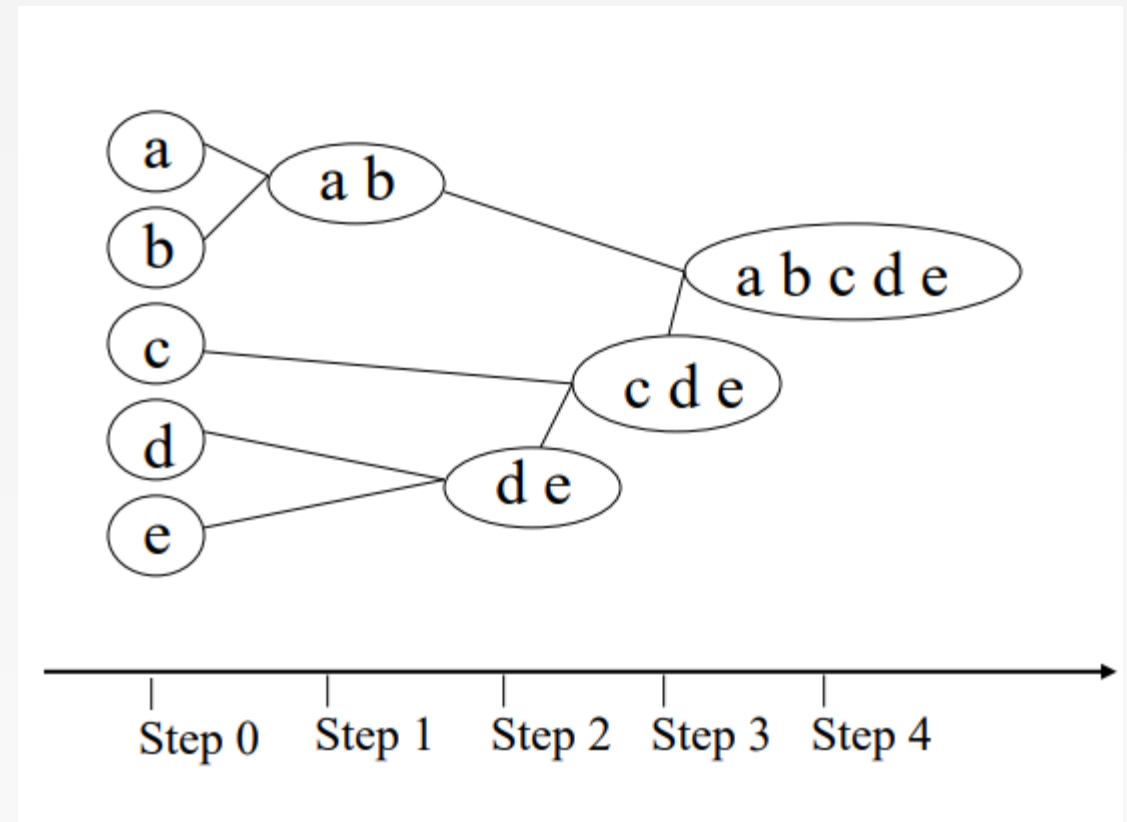
Types

- Divisive (top-down) clustering
 - All objects in one cluster
 - Select a cluster and split it into two sub clusters
 - Until each leaf cluster contains only one object



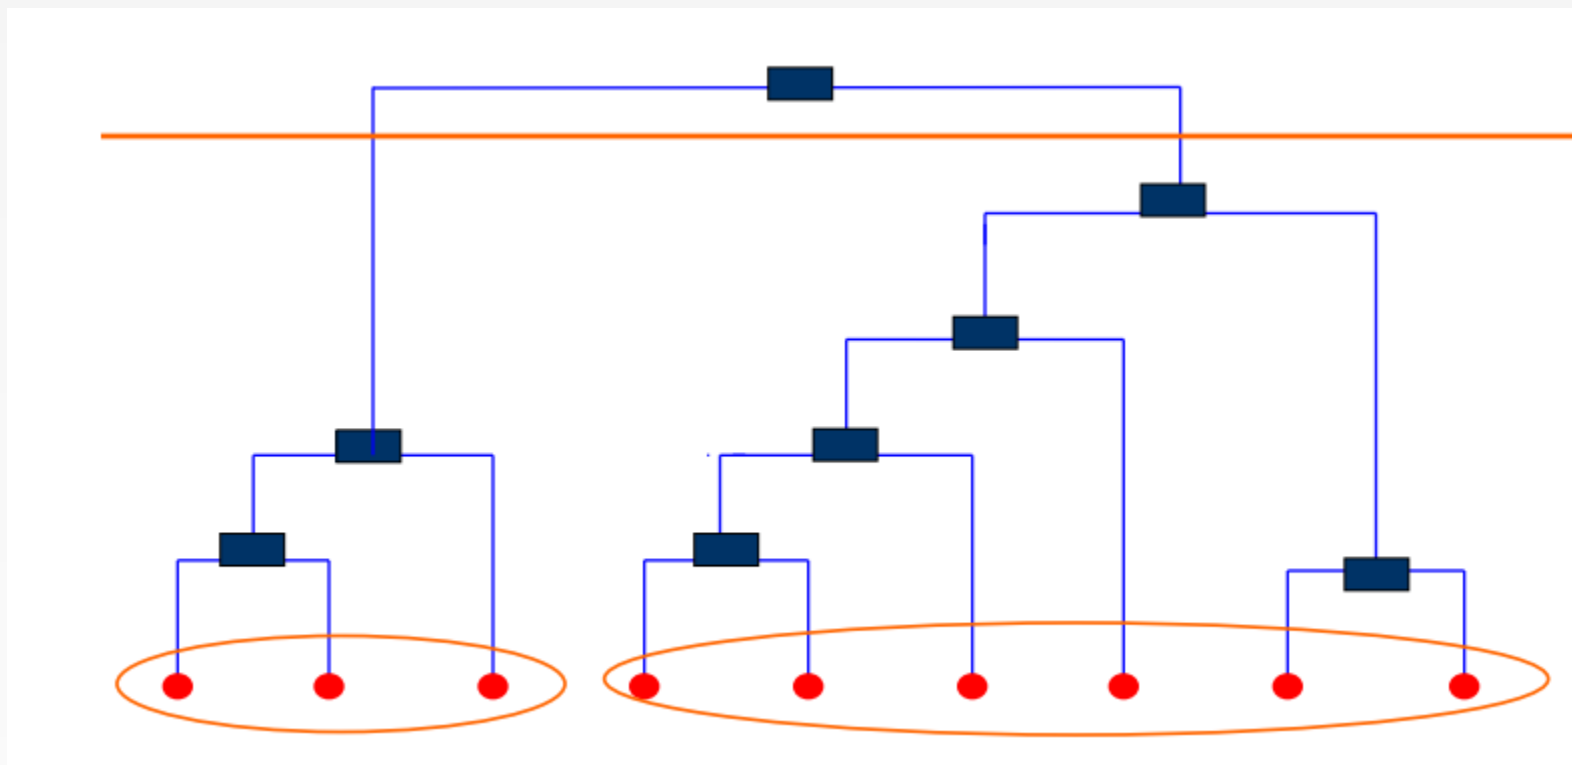
Types

- Agglomerative (bottom-up) clustering
 - Each object is a cluster
 - Merge two clusters which are most similar to each other
 - Until all objects are merged into a single cluster




Dendrogram

- A tree that shows how clusters are merged/split hierarchically
- Each node on the tree is a cluster; each leaf node is a singleton cluster
- A clustering of the data objects is obtained by cutting the dendrogram at the desired level, then each connected component forms a cluster




Inter-Cluster Distance




MIN (Single Link)

- The distance between two clusters is represented by the distance of the closest pair of data objects belonging to different clusters.
- Determined by one pair of points, i.e., by one link in the proximity graph




- Limitation: sensitive to noise/outliers




Original Points

- **Sensitive to noise and outliers**




Two Clusters



Three Clusters



MAX (Complete link)

- The distance between two clusters is represented by the distance of the farthest pair of data objects belonging to different clusters




MAX (Complete link)

- Strength: less sensitive to noise/outliers
- Limitations: tends to break large clusters




Group average

- The distance between two clusters is represented by the average distance of all pairs of data objects belonging to different clusters
- Determined by all pairs of points in the two clusters



Centroid Distance

- The distance between two clusters is represented by the distance between the centers of the clusters
- – Determined by cluster centroids



Ward's Method

- Similarity of two clusters is based on the increase in squared error when two clusters are merged
- Similar to group average if distance between points is distance squared
- Less susceptible to noise and outliers