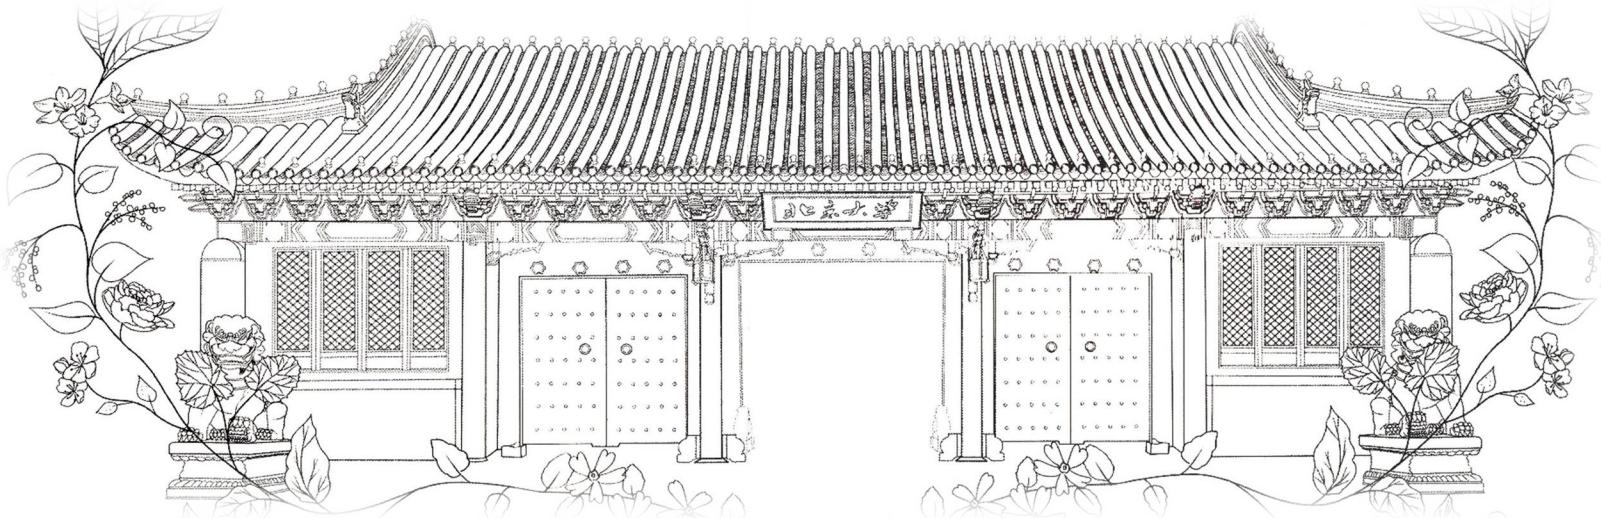


《Python数据分析》

应用篇：时间序列分析基础





三种常见数据

- 时间序列数据（Time series data），是一批按照时间先后顺序排列的统计数据
- 截面数据（Cross-section data），是一批发生在同一时间截面上的数据
 - 工业普查数据、人口普查数据、家庭调查数据等
- 面板数据（Panel data），是时间序列数据与截面数据的合成体。
 - 1978-1999年我国各省市城镇居民消费结构的调查资料



三种常见数据

时间序列数据
截面数据
面板数据

表1 华东地区各省市 GDP 历史数据 单位：亿元

	1995	1996	1997	1998	1999
上海	2462.57	2902.20	3360.21	3688.20	4034.96
江苏	5155.25	6004.21	6680.34	7199.95	7697.82
浙江	3524.79	4146.06	4638.24	4987.50	5364.89
安徽	2003.66	2339.25	2669.95	2805.45	2908.59
福建	2191.27	2583.83	3000.36	3286.56	3550.24
江西	1244.04	1517.26	1715.18	1851.98	1962.98
山东	4996.87	5960.42	6650.02	7162.20	7662.10

数据来源：中国统计年鉴 1996-2000。

北京交通大学 HD:72085105





时间序列数据

- 形式上由现象所属的时间和现象在不同时间上的观察值两部分组成。
- 时间序列的时间是变化的。常用的时间间隔有：年、季度、月、周、日……
- 时间序列数据有时存在季节变动和序列相关——自相关。

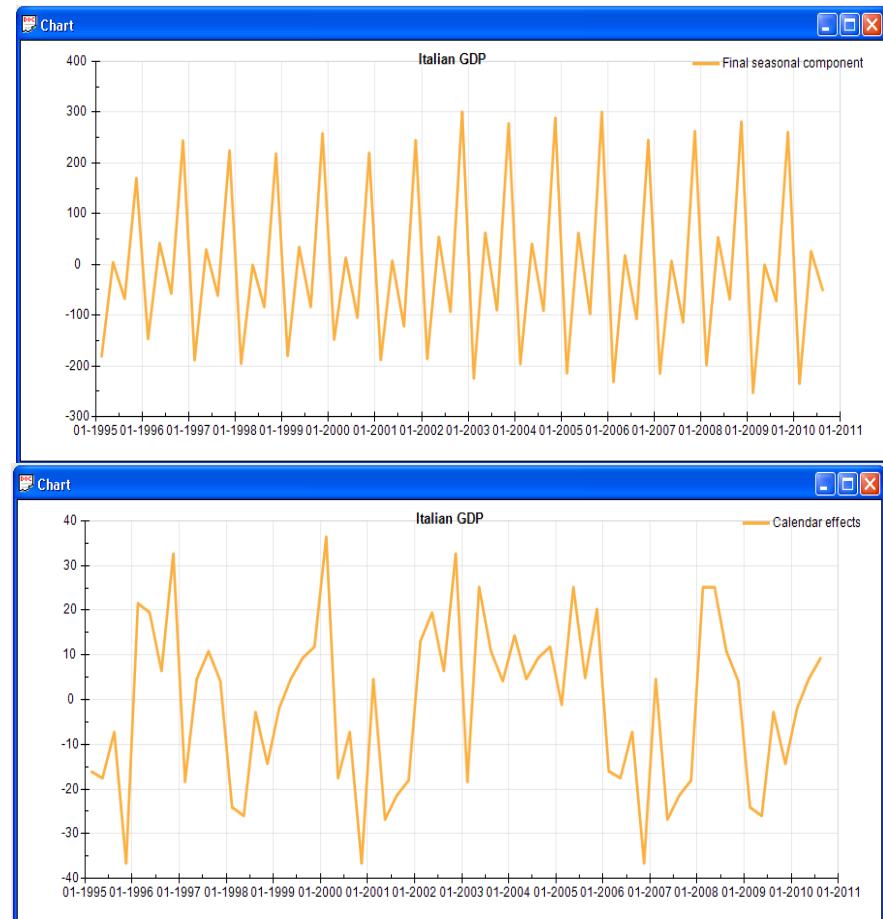




时间序列

- 季度性时间序列

- 日期性时间序列





时间序列

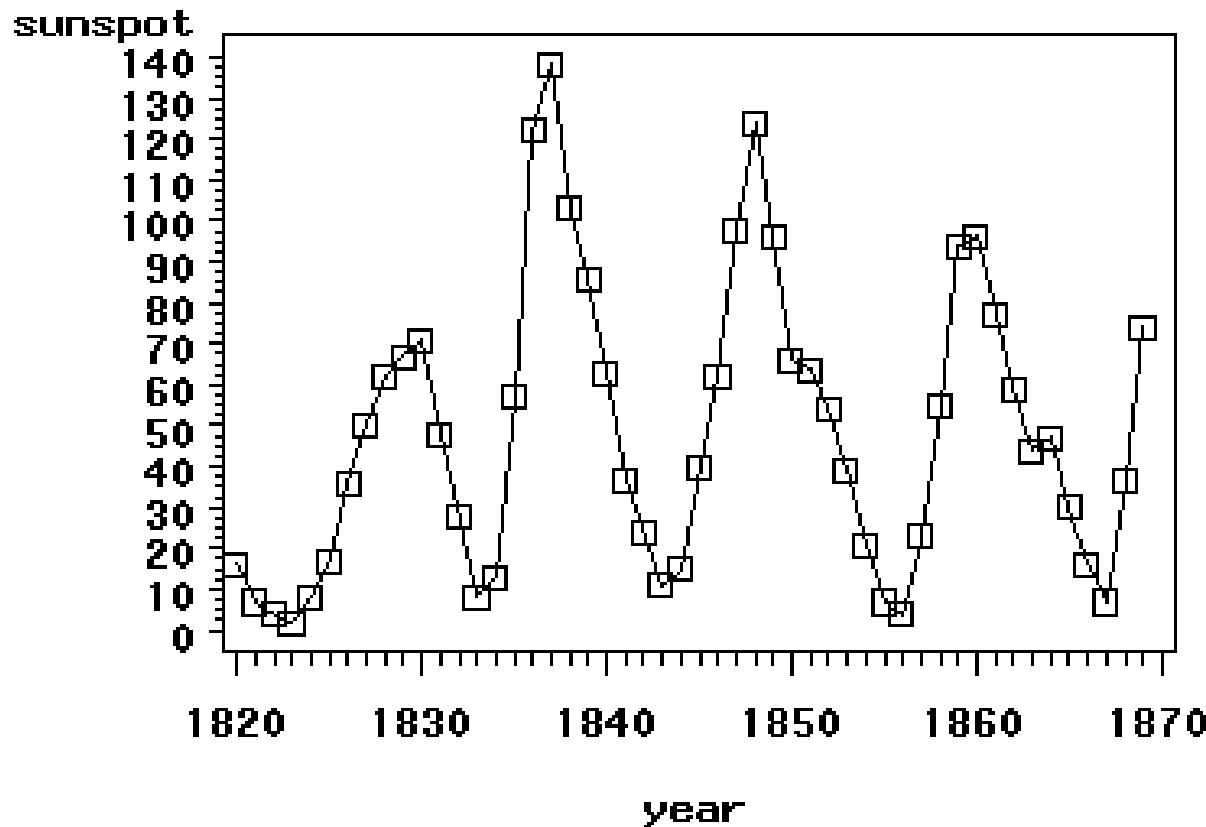
- 一个时间序列是依时间顺序生成的观测值的集合
 - 若该集合是连续的 \rightarrow 连续型时间序列
 - 若该集合是离散的 \rightarrow 离散型时间序列
- 设 y_t 是时间序列在时刻（或时期） t 的观测值，当在 $t = 1, 2, 3, 4, 5, 6, 7, \dots, n$ 采样时，得到时间序列：

$$y_1, y_2, y_3, y_4, y_5, y_6, \dots, y_n$$



时间序列

- 德国业余天文学家施瓦尔发现太阳黑子的活动具有11年左右的周期





目录

- Python 中时间和日期的处理
- 时间序列基础
- 日期的范围、频率和转化
- 时间区间和区间算术
- 时间序列的重采样 (resampling)
- 移动窗口函数 (moving window functions)





Python中时间和日期的处理





常用方法和函数？

```
In [10]: from datetime import datetime
```

```
In [11]: now = datetime.now()
```

1秒=1000000 微秒(μs)

```
In [12]: now
```

```
Out[12]: datetime.datetime(2017, 9, 25, 14, 5, 52, 72973)
```

```
In [13]: now.year, now.month, now.day
```

```
Out[13]: (2017, 9, 25)
```





常用方法和函数

```
In [14]: delta = datetime(2011, 1, 7) - datetime(2008, 6, 24, 8, 15)
```

```
In [15]: delta
```

```
Out[15]: datetime.timedelta(926, 56700)
```

```
In [16]: delta.days
```

```
Out[16]: 926
```

```
In [17]: delta.seconds
```

```
Out[17]: 56700
```

```
In [18]: from datetime import timedelta
```

```
In [19]: start = datetime(2011, 1, 7)
```

```
In [20]: start + timedelta(12)
```

```
Out[20]: datetime.datetime(2011, 1, 19, 0, 0)
```

```
In [21]: start - 2 * timedelta(12)
```

```
Out[21]: datetime.datetime(2010, 12, 14, 0, 0)
```

```
now
```

```
datetime.datetime(2018, 1, 20, 11, 14, 45, 545000)
```

```
now-datetime.timedelta(hours=3, minutes=30)
```

```
datetime.datetime(2018, 1, 20, 7, 44, 45, 545000)
```





常见的几种数据类型

公历

Type	Description
date	Store calendar date (year, month, day) using the Gregorian calendar
time	Store time of day as hours, minutes, seconds, and microseconds
datetime	Stores both date and time
timedelta	Represents the difference between two datetime values (as days, seconds, and microseconds)
tzinfo	Base type for storing time zone information





字符串和日期之间的转化

```
In [22]: stamp = datetime(2011, 1, 3)
```

```
In [23]: str(stamp)
```

```
Out[23]: '2011-01-03 00:00:00'
```

```
In [24]: stamp.strftime('%Y-%m-%d')
```

```
Out[24]: '2011-01-03'
```

Type	Description
%Y	Four-digit year
%y	Two-digit year
%m	Two-digit month [01, 12]
%d	Two-digit day [01, 31]
%H	Hour (24-hour clock) [00, 23]
%I	Hour (12-hour clock) [01, 12]
%M	Two-digit minute [00, 59]
%S	Second [00, 61] (seconds 60, 61 account for leap seconds)
%w	Weekday as integer [0 (Sunday), 6]
%U	Week number of the year [00, 53]; Sunday is considered the first day of the week, and days before the first Sunday of the year are "week 0"
%W	Week number of the year [00, 53]; Monday is considered the first day of the week, and days before the first Monday of the year are "week 0"
%z	UTC time zone offset as +HHMM or -HHMM; empty if time zone naive
%F	Shortcut for %Y-%m-%d (e.g., 2012-4-18)
%D	Shortcut for %m/%d/%y (e.g., 04/18/12)

60、61算作闰秒





字符串和日期之间的转化

```
In [25]: value = '2011-01-03'
```

```
In [26]: datetime.strptime(value, '%Y-%m-%d')  
Out[26]: datetime.datetime(2011, 1, 3, 0, 0)
```

```
In [27]: datestrs = ['7/6/2011', '8/6/2011']
```

```
In [28]: [datetime.strptime(x, '%m/%d/%Y') for x in datestrs]
```

```
Out[28]:  
[datetime.datetime(2011, 7, 6, 0, 0),  
 datetime.datetime(2011, 8, 6, 0, 0)]
```





字符串和日期之间的转化

- Datetime.strptime在已知格式时转换日期很好用，但每次都要编写格式代码可能很不方便
- 使用dateutil包的parser.parse方法

```
In [29]: from dateutil.parser import parse
```

```
In [30]: parse('2011-01-03')
Out[30]: datetime.datetime(2011, 1, 3, 0, 0)
```

dateutil is capable of parsing most human-intelligible date representations:

```
In [31]: parse('Jan 31, 1997 10:45 PM')
Out[31]: datetime.datetime(1997, 1, 31, 22, 45)
```

In international locales, day appearing before month is very common, so you can pass dayfirst=True to indicate this:

```
In [32]: parse('6/12/2011', dayfirst=True)
Out[32]: datetime.datetime(2011, 12, 6, 0, 0)
```





字符串和日期之间的转化: pandas

```
In [33]: datestrs = ['2011-07-06 12:00:00', '2011-08-06 00:00:00']
```

```
In [34]: pd.to_datetime(datestrs)
```

```
Out[34]: DatetimeIndex(['2011-07-06 12:00:00', '2011-08-06 00:00:00'], dtype='datetime64[ns]', freq=None)
```

It also handles values that should be considered missing (None, empty string, etc.):

```
In [35]: idx = pd.to_datetime(datestrs + [None])
```

```
In [36]: idx
```

```
Out[36]: DatetimeIndex(['2011-07-06 12:00:00', '2011-08-06 00:00:00', 'NaT'], dtype='datetime64[ns]', freq=None)
```

```
In [37]: idx[2]
```

```
Out[37]: NaT
```

Not a time

```
In [38]: pd.isnull(idx)
```

```
Out[38]: array([False, False, True], dtype=bool)
```





时间序列基础





pandas 中的时间序列

- Pandas 中的基础时间序列种类是由时间戳索引的 Series，在 pandas 外部则通常表示为 Python 字符串 或 datetime 对象。

```
In [39]: from datetime import datetime
```

```
In [40]: dates = [datetime(2011, 1, 2), datetime(2011, 1, 5),
.....           datetime(2011, 1, 7), datetime(2011, 1, 8),
.....           datetime(2011, 1, 10), datetime(2011, 1, 12)]
```

```
In [41]: ts = pd.Series(np.random.randn(6), index=dates)
```

```
In [42]: ts
Out[42]:
2011-01-02    -0.204708
2011-01-05     0.478943
2011-01-07    -0.519439
2011-01-08    -0.555730
2011-01-10     1.965781
2011-01-12     1.393406
dtype: float64
```

```
In [43]: ts.index
Out[43]:
DatetimeIndex(['2011-01-02', '2011-01-05', '2011-01-07', '2011-01-08',
               '2011-01-10', '2011-01-12'],
              dtype='datetime64[ns]', freq=None)
```





pandas 中的时间序列

- 和其他Series类似，不同索引的时间序列之间的算术运算在日期上自动对齐。

```
In [44]: ts + ts[::2]
```

```
Out[44]:
```

```
2011-01-02    -0.409415
2011-01-05        NaN
2011-01-07    -1.038877
2011-01-08        NaN
2011-01-10    3.931561
2011-01-12        NaN
dtype: float64
```

- Pandas使用Numpy的datetime64数据类型在纳米级分辨率下存储的时间戳：

```
In [45]: ts.index.dtype
```

```
Out[45]: dtype('<M8[ns]')
```

```
In [42]: ts
Out[42]:
2011-01-02    -0.204708
2011-01-05    0.478943
2011-01-07   -0.519439
2011-01-08   -0.555730
2011-01-10    1.965781
2011-01-12    1.393406
dtype: float64
```





pandas中的时间序列

- DatetimeIndex的标量值是pandas的Timestamp对象。

```
In [46]: stamp = ts.index[0]
```

```
In [47]: stamp  
Out[47]: Timestamp('2011-01-02 00:00:00')
```

- 所有使用datetime对象的地方都可以使用Timestamp。
 - Timestamp还可以存储频率信息（如有）并了解如何进行时区转换和其他类型操作。

```
In [42]: ts  
Out[42]:  
2011-01-02    -0.204708  
2011-01-05     0.478943  
2011-01-07    -0.519439  
2011-01-08    -0.555730  
2011-01-10     1.965781  
2011-01-12     1.393406  
dtype: float64
```





索引、选择、子序列

- 当基于标签进行索引和选择时，时间序列的行为和其他的pandas.Series类似：

```
In [48]: stamp = ts.index[2]
```

```
In [49]: ts[stamp]  
Out[49]: -0.51943871505673811
```

- 为了方便，还可以传递一个能解释为日期的字符串：

```
In [50]: ts['1/10/2011']  
Out[50]: 1.9657805725027142
```

```
In [51]: ts['20110110']  
Out[51]: 1.9657805725027142
```

```
In [42]: ts  
Out[42]:  
2011-01-02    -0.204708  
2011-01-05     0.478943  
2011-01-07    -0.519439  
2011-01-08    -0.555730  
2011-01-10     1.965781  
2011-01-12     1.393406  
dtvde: float64
```





索引、选择、子序列

- 对一个长的时间序列，可以传递一个年份或一个年份和 periods 值来选择数据的切片：

```
In [52]: longer_ts = pd.Series(np.random.randn(1000),
....:                         index=pd.date_range('1/1/2000', periods=1000))
```

```
In [53]: longer_ts
Out[53]:
2000-01-01    0.092908  2002-09-22    0.930944
2000-01-02    0.281746  2002-09-23   -0.811676
2000-01-03    0.769023  2002-09-24   -1.830156
2000-01-04    1.246435  2002-09-25   -0.138730
2000-01-05    1.007189  2002-09-26    0.334088
2000-01-06   -1.296221
2000-01-07    0.274992
2000-01-08    0.228913
2000-01-09    1.352917
2000-01-10    0.886429
...
2002-09-17   -0.139298
2002-09-18   -1.159926
2002-09-19    0.618965
2002-09-20    1.373890
2002-09-21   -0.983505
```

```
In [54]: longer_ts['2001']
Out[54]:
2001-01-01    1.599534
2001-01-02    0.474071
2001-01-03    0.151326
2001-01-04   -0.542173
2001-01-05   -0.475496
2001-01-06    0.106403
2001-01-07   -1.308228
2001-01-08    2.173185
2001-01-09    0.564561
2001-01-10   -0.190481
...
2001-12-22    0.000369
2001-12-23    0.900885
2001-12-24   -0.454869
2001-12-25   -0.864547
2001-12-26    1.129120
2001-12-27    0.057874
2001-12-28   -0.433739
2001-12-29    0.092698
2001-12-30   -1.397820
2001-12-31    1.457823
Freq: D, Length: 365, dtype: float64
```



索引、选择、子序列

Here, the string '2001' is interpreted as a year and selects that time period. This also works if you specify the month:

```
In [55]: longer_ts['2001-05']  
Out[55]:  
2001-05-01    -0.622547    2001-05-27    0.235477  
2001-05-02     0.936289    2001-05-28    0.111835  
2001-05-03     0.750018    2001-05-29   -1.251504  
2001-05-04    -0.056715    2001-05-30   -2.949343  
2001-05-05     2.300675    2001-05-31    0.634634  
2001-05-06     0.569497  
2001-05-07     1.489410  
2001-05-08     1.264250  
2001-05-09    -0.761837  
2001-05-10    -0.331617  
...  
2001-05-22     0.503699  
2001-05-23    -1.387874  
2001-05-24     0.204851  
2001-05-25     0.603705  
2001-05-26     0.545680  
Freq: D, Length: 31, dtype: float64
```





索引、选择、子序列

- 也可以使用datetime对象进行切片：

```
In [56]: ts[datetime(2011, 1, 7):]
```

```
Out[56]:
```

```
2011-01-07    -0.519439
```

```
2011-01-08    -0.555730
```

```
2011-01-10    1.965781
```

```
2011-01-12    1.393406
```

```
dtype: float64
```

- 由于大多数时间序列数据是按时间顺序排序的，可以使用不包含在时间序列中的时间戳进行切片，以执行范围查询：

```
In [58]: ts['1/6/2011':'1/11/2011']
```

```
Out[58]:
```

```
2011-01-07    -0.519439
```

```
2011-01-08    -0.555730
```

```
2011-01-10    1.965781
```

```
dtype: float64
```

```
In [42]: ts
Out[42]:
2011-01-02    -0.204708
2011-01-05    0.478943
2011-01-07    -0.519439
2011-01-08    -0.555730
2011-01-10    1.965781
2011-01-12    1.393406
dtype: float64
```





索引、选择、子序列

- 有一个等价实例方法，truncate，它可以在两个日期间对 Series 进行切片：

```
In [59]: ts.truncate(after='1/9/2011')
Out[59]:
2011-01-02    -0.204708
2011-01-05     0.478943
2011-01-07    -0.519439
2011-01-08    -0.555730
dtype: float64
```

```
In [42]: ts
Out[42]:
2011-01-02    -0.204708
2011-01-05     0.478943
2011-01-07    -0.519439
2011-01-08    -0.555730
2011-01-10     1.965781
2011-01-12     1.393406
dtype: float64
```





索引、选择、子序列

- 以上的很多操作也都适用于DataFrame，并在其行上进行索引：

```
In [60]: dates = pd.date_range('1/1/2000', periods=100, freq='W-WED')
```

```
In [61]: long_df = pd.DataFrame(np.random.randn(100, 4),
.....:                               index=dates,
.....:                               columns=['Colorado', 'Texas',
.....:                                         'New York', 'Ohio'])
```

```
In [62]: long_df.loc['5-2001']
```

```
Out[62]:
```

	Colorado	Texas	New York	Ohio
2001-05-02	-0.006045	0.490094	-0.277186	-0.707213
2001-05-09	-0.560107	2.735527	0.927335	1.513906
2001-05-16	0.538600	1.273768	0.667876	-0.969206
2001-05-23	1.676091	-0.817649	0.050188	1.951312
2001-05-30	3.260383	0.963301	1.201206	-1.852001





【补充】 freq参数

Alias	Description
B	business day frequency
C	custom business day frequency
D	calendar day frequency
W	weekly frequency
M	month end frequency
SM	semi-month end frequency (15th and end of month)
BM	business month end frequency
CBM	custom business month end frequency
MS	month start frequency
SMS	semi-month start frequency (1st and 15th)
BMS	business month start frequency
CBMS	custom business month start frequency
Q	quarter end frequency
BQ	business quarter end frequency
QS	quarter start frequency
BQS	business quarter start frequency
A, Y	year end frequency
BA, BY	business year end frequency
AS, YS	year start frequency
BAS, BYS	business year start frequency
BH	business hour frequency
H	hourly frequency
T, min	minutely frequency
S	secondly frequency
L, ms	milliseconds
U, us	microseconds
N	nanoseconds



含有重复索引的时间序列

```
In [63]: dates = pd.DatetimeIndex(['1/1/2000', '1/2/2000', '1/2/2000',
.....                           '1/2/2000', '1/3/2000'])
```

```
In [64]: dup_ts = pd.Series(np.arange(5), index=dates)
```

```
In [65]: dup_ts
```

```
Out[65]:
```

```
2000-01-01    0
2000-01-02    1
2000-01-02    2
2000-01-02    3
2000-01-03    4
dtype: int64
```

We can tell that the index is not unique by checking its `is_unique` property:

```
In [66]: dup_ts.index.is_unique
Out[66]: False
```





```
In [65]: dup_ts  
Out[65]:  
2000-01-01    0  
2000-01-02    1  
2000-01-02    2  
2000-01-02    3  
2000-01-03    4
```

含有重复索引的时间序列

- 对上面的Series进行索引，结果是标量值还是Series切片取决于是否有时间戳是重复的：

```
In [67]: dup_ts['1/3/2000'] # not duplicated  
Out[67]: 4
```

```
In [68]: dup_ts['1/2/2000'] # duplicated  
Out[68]:  
2000-01-02    1  
2000-01-02    2  
2000-01-02    3  
dtype: int64
```



含有重复索引的时间序列

- 假设我们想要聚合含有非唯一时间戳的数据，一种方式就是使用groupby并传递level = 0：

```
In [69]: grouped = dup_ts.groupby(level=0)
```

```
In [70]: grouped.mean()
```

```
Out[70]:
```

```
2000-01-01    0
```

```
2000-01-02    2
```

```
2000-01-03    4
```

```
dtype: int64
```

```
In [71]: grouped.count()
```

```
Out[71]:
```

```
2000-01-01    1
```

```
2000-01-02    3
```

```
2000-01-03    1
```

```
dtype: int64
```

```
In [65]: dup_ts
```

```
Out[65]:
```

2000-01-01	0
2000-01-02	1
2000-01-02	2
2000-01-02	3
2000-01-03	4

可以用在什么场景呢？





日期的范围、频率和转化





日期的范围

- 再次强调：Pandas.date_range是用于根据特定频率生成指定长度的DateTimeIndex
 - 默认情况下，date_range生成的是每日的时间戳

```
In [74]: index = pd.date_range('2012-04-01', '2012-06-01')
```

```
In [75]: index
```

```
Out[75]:
```

```
DatetimeIndex(['2012-04-01', '2012-04-02', '2012-04-03', '2012-04-04',
                 '2012-04-05', '2012-04-06', '2012-04-07', '2012-04-08',
                 '2012-04-09', '2012-04-10', '2012-04-11', '2012-04-12',
                 '2012-04-13', '2012-04-14', '2012-04-15', '2012-04-16',
                 '2012-04-17', '2012-04-18', '2012-04-19', '2012-04-20',
                 '2012-04-21', '2012-04-22', '2012-04-23', '2012-04-24',
                 '2012-04-25', '2012-04-26', '2012-04-27', '2012-04-28',
                 '2012-04-29', '2012-04-30', '2012-05-01', '2012-05-02',
                 '2012-05-03', '2012-05-04', '2012-05-05', '2012-05-06',
                 '2012-05-07', '2012-05-08', '2012-05-09', '2012-05-10',
                 '2012-05-11', '2012-05-12', '2012-05-13', '2012-05-14',
                 '2012-05-15', '2012-05-16', '2012-05-17', '2012-05-18',
                 '2012-05-19', '2012-05-20', '2012-05-21', '2012-05-22',
                 '2012-05-23', '2012-05-24', '2012-05-25', '2012-05-26',
                 '2012-05-27', '2012-05-28', '2012-05-29', '2012-05-30',
                 '2012-05-31', '2012-06-01'],
                dtype='datetime64[ns]', freq='D')
```





日期的范围

- 如果只传递一个起始或结尾日期，必须传递一个用于生成范围的数字：

```
In [76]: pd.date_range(start='2012-04-01', periods=20)
```

```
Out[76]:
```

```
DatetimeIndex(['2012-04-01', '2012-04-02', '2012-04-03', '2012-04-04',
                 '2012-04-05', '2012-04-06', '2012-04-07', '2012-04-08',
                 '2012-04-09', '2012-04-10', '2012-04-11', '2012-04-12',
                 '2012-04-13', '2012-04-14', '2012-04-15', '2012-04-16',
                 '2012-04-17', '2012-04-18', '2012-04-19', '2012-04-20'],
                dtype='datetime64[ns]', freq='D')
```

```
In [77]: pd.date_range(end='2012-06-01', periods=20)
```

```
Out[77]:
```

```
DatetimeIndex(['2012-05-13', '2012-05-14', '2012-05-15', '2012-05-16',
                 '2012-05-17', '2012-05-18', '2012-05-19', '2012-05-20',
                 '2012-05-21', '2012-05-22', '2012-05-23', '2012-05-24',
                 '2012-05-25', '2012-05-26', '2012-05-27', '2012-05-28',
                 '2012-05-29', '2012-05-30', '2012-05-31', '2012-06-01'],
                dtype='datetime64[ns]', freq='D')
```



日期的范围

BusinessMonthEnd, 工作日的月底日期

```
In [78]: pd.date_range('2000-01-01', '2000-12-01', freq='BM')
```

```
Out[78]:
```

```
DatetimeIndex(['2000-01-31', '2000-02-29', '2000-03-31', '2000-04-28',
                 '2000-05-31', '2000-06-30', '2000-07-31', '2000-08-31',
                 '2000-09-29', '2000-10-31', '2000-11-30'],
                dtype='datetime64[ns]', freq='BM')
```



日期的范围

- 有时候你会获得包含时间信息的开始日期或结束日期，但是你想要生成的是标准化为0:00的时间戳。此时可使用normalize选项实现这个功能：

```
In [80]: pd.date_range('2012-05-02 12:56:31', periods=5, normalize=True)
```

```
Out[80]:
```

```
DatetimeIndex(['2012-05-02', '2012-05-03', '2012-05-04', '2012-05-05',
                 '2012-05-06'],
                dtype='datetime64[ns]', freq='D')
```

没有normalize的时候

```
1 pd.date_range('2012-05-02 12:56:31', periods=5)
```

```
DatetimeIndex(['2012-05-02 12:56:31', '2012-05-03 12:56:31',
                 '2012-05-04 12:56:31', '2012-05-05 12:56:31',
                 '2012-05-06 12:56:31'],
                dtype='datetime64[ns]', freq='D')
```





日期的频率

- Pandas 中的频率是由基础频率和倍数组成的。
 - 基础频率会有字符串别名，如 M 表示每月，H 表示每小时
 - 对每个基础频率，都有一个对象可以被用于定义日期偏置：

```
In [81]: from pandas.tseries.offsets import Hour, Minute
```

```
In [82]: hour = Hour()
```

```
In [83]: hour
```

```
Out[83]: <Hour>
```

```
In [86]: pd.date_range('2000-01-01', '2000-01-03 23:59', freq='4h')
```

```
Out[86]:
```

```
DatetimeIndex(['2000-01-01 00:00:00', '2000-01-01 04:00:00',
                 '2000-01-01 08:00:00', '2000-01-01 12:00:00',
                 '2000-01-01 16:00:00', '2000-01-01 20:00:00',
                 '2000-01-02 00:00:00', '2000-01-02 04:00:00',
                 '2000-01-02 08:00:00', '2000-01-02 12:00:00',
                 '2000-01-02 16:00:00', '2000-01-02 20:00:00',
                 '2000-01-03 00:00:00', '2000-01-03 04:00:00',
                 '2000-01-03 08:00:00', '2000-01-03 12:00:00',
                 '2000-01-03 16:00:00', '2000-01-03 20:00:00'],
                dtype='datetime64[ns]', freq='4H')
```





日期的频率

```
In [87]: Hour(2) + Minute(30)
```

```
Out[87]: <150 * Minutes>
```

```
In [88]: pd.date_range('2000-01-01', periods=10, freq='1h30min')
```

```
Out[88]:
```

```
DatetimeIndex(['2000-01-01 00:00:00', '2000-01-01 01:30:00',
                 '2000-01-01 03:00:00', '2000-01-01 04:30:00',
                 '2000-01-01 06:00:00', '2000-01-01 07:30:00',
                 '2000-01-01 09:00:00', '2000-01-01 10:30:00',
                 '2000-01-01 12:00:00', '2000-01-01 13:30:00'],
                dtype='datetime64[ns]', freq='90T')
```





日期的频率：补充一个例子

```
1 rng = pd.date_range('2012-01-01', '2012-09-01', freq='WOM-3FRI')  
2 list(rng)
```

```
[Timestamp('2012-01-20 00:00:00', freq='WOM-3FRI'),  
 Timestamp('2012-02-17 00:00:00', freq='WOM-3FRI'),  
 Timestamp('2012-03-16 00:00:00', freq='WOM-3FRI'),  
 Timestamp('2012-04-20 00:00:00', freq='WOM-3FRI'),  
 Timestamp('2012-05-18 00:00:00', freq='WOM-3FRI'),  
 Timestamp('2012-06-15 00:00:00', freq='WOM-3FRI'),  
 Timestamp('2012-07-20 00:00:00', freq='WOM-3FRI'),  
 Timestamp('2012-08-17 00:00:00', freq='WOM-3FRI')]
```

“月中某星期”
(week of month) 表
示为WOM，如每月第
三个星期五





日期的“移位”

- “移位”是指将日期按时间向前移动或向后移动。
 - Series和DataFrame都有一个shift方法用于进行简单的前、后向移位，而不改变索引。

```
In [91]: ts = pd.Series(np.random.randn(4),  
.....: index=pd.date_range('1/1/2000', periods=4, freq='M'))
```

```
In [92]: ts  
Out[92]:  
2000-01-31    -0.066748  
2000-02-29     0.838639  
2000-03-31    -0.117388  
2000-04-30    -0.517795  
Freq: M, dtype: float64
```

```
In [93]: ts.shift(2)  
Out[93]:  
2000-01-31      NaN  
2000-02-29      NaN  
2000-03-31    -0.066748  
2000-04-30     0.838639  
Freq: M, dtype: float64
```

```
In [94]: ts.shift(-2)  
Out[94]:  
2000-01-31    -0.117388  
2000-02-29    -0.517795  
2000-03-31      NaN  
2000-04-30      NaN  
Freq: M, dtype: float64
```





日期的“移位”

If freq is specified then the index values are shifted but the data is not realigned.

In [95]: ts.shift(2, freq='M')
Out[95]:

2000-03-31 -0.066748
2000-04-30 0.838639
2000-05-31 -0.117388
2000-06-30 -0.517795

Freq: M, dtype: float64

In [92]: ts
Out[92]:
2000-01-31 -0.066748
2000-02-29 0.838639
2000-03-31 -0.117388
2000-04-30 -0.517795

In [96]: ts.shift(3, freq='D')
Out[96]:
2000-02-03 -0.066748
2000-03-03 0.838639
2000-04-03 -0.117388
2000-05-03 -0.517795
dtype: float64

In [97]: ts.shift(1, freq='90T')
Out[97]:
2000-01-31 01:30:00 -0.066748
2000-02-29 01:30:00 0.838639
2000-03-31 01:30:00 -0.117388
2000-04-30 01:30:00 -0.517795
Freq: M, dtype: float64

The T here stands for minutes.



日期的“移位”

- Pandas日期偏置也可以使用datetime或Timestamp对象完成

```
In [98]: from pandas.tseries.offsets import Day, MonthEnd
```

```
In [99]: now = datetime(2011, 11, 17)
```

```
In [100]: now + 3 * Day()
```

```
Out[100]: Timestamp('2011-11-20 00:00:00')
```

```
In [101]: now + MonthEnd()
```

```
Out[101]: Timestamp('2011-11-30 00:00:00')
```

```
In [102]: now + MonthEnd(2)
```

```
Out[102]: Timestamp('2011-12-31 00:00:00')
```





日期的“移位”

```
In [99]: now = datetime(2011, 11, 17)
```

```
In [103]: offset = MonthEnd()
```

```
In [104]: offset.rollforward(now)
```

```
Out[104]: Timestamp('2011-11-30 00:00:00')
```

```
In [105]: offset.rollback(now)
```

```
Out[105]: Timestamp('2011-10-31 00:00:00')
```





日期的“移位”

```
In [106]: ts = pd.Series(np.random.randn(20),  
..... index=pd.date_range('1/15/2000', periods=20, freq='4d'))
```

```
In [107]: ts
```

```
Out[107]:  
2000-01-15    -0.116696  
2000-01-19     2.389645  
2000-01-23    -0.932454  
2000-01-27    -0.229331  
2000-01-31    -1.140330  
2000-02-04     0.439920  
2000-02-08    -0.823758  
2000-02-12    -0.520930  
2000-02-16     0.350282  
2000-02-20     0.204395  
2000-02-24     0.133445  
2000-02-28     0.327905  
2000-03-03     0.072153  
2000-03-07     0.131678  
2000-03-11    -1.297459  
2000-03-15     0.997747  
2000-03-19     0.870955  
2000-03-23    -0.991253  
2000-03-27     0.151699  
2000-03-31     1.266151  
Freq: 4D, dtype: float64
```

```
In [108]: ts.groupby(offset.rollforward).mean()
```

```
Out[108]:  
2000-01-31    -0.005833  
2000-02-29     0.015894  
2000-03-31     0.150209  
dtype: float64
```





时区

```
In [110]: import pytz
```

```
In [111]: pytz.common_timezones[-5:]
```

```
Out[111]: ['US/Eastern', 'US/Hawaii', 'US/Mountain', 'US/Pacific', 'UTC']
```

To get a time zone object from pytz, use `pytz.timezone`:

```
In [112]: tz = pytz.timezone('America/New_York')
```

```
In [113]: tz
```

```
Out[113]: <DstTzInfo 'America/New_York' LMT-1 day, 19:04:00 STD>
```

Methods in pandas will accept either time zone names or these objects.

Time Zone Localization and Conversion

By default, time series in pandas are *time zone naive*. For example, consider the following time series:

```
In [114]: rng = pd.date_range('3/9/2012 9:30', periods=6, freq='D')
```

```
In [115]: ts = pd.Series(np.random.randn(len(rng)), index=rng)
```

```
In [116]: ts
```

```
Out[116]:
```

```
2012-03-09 09:30:00    -0.202469
2012-03-10 09:30:00     0.050718
2012-03-11 09:30:00     0.639869
2012-03-12 09:30:00     0.597594
```

```
2012-03-13 09:30:00    -0.797246
2012-03-14 09:30:00     0.472879
Freq: D, dtype: float64
```

The index's `tz` field is `None`:

```
In [117]: print(ts.index.tz)
None
```

Date ranges can be generated with a time zone set:

```
In [118]: pd.date_range('3/9/2012 9:30', periods=10, freq='D', tz='UTC')
Out[118]:
DatetimeIndex(['2012-03-09 09:30:00+00:00', '2012-03-10 09:30:00+00:00',
               '2012-03-11 09:30:00+00:00', '2012-03-12 09:30:00+00:00',
               '2012-03-13 09:30:00+00:00', '2012-03-14 09:30:00+00:00',
               '2012-03-15 09:30:00+00:00', '2012-03-16 09:30:00+00:00',
               '2012-03-17 09:30:00+00:00', '2012-03-18 09:30:00+00:00'],
              'datetime64[ns, UTC]', freq='D')
```





时间区间和区间算术





时间区间

- 时间区间表示的是时间范围，使用Period类：

```
In [149]: p = pd.Period(2007, freq='A-DEC')
```

给定月份（DEC）所在月的最后一个工作日所对应的年度日期

```
In [150]: p
```

```
Out[150]: Period('2007', 'A-DEC')
```

两个区间频率相同，差值是它们间的单位数

```
In [151]: p + 5
```

```
Out[151]: Period('2012', 'A-DEC')
```

```
In [152]: p - 2
```

```
Out[152]: Period('2005', 'A-DEC')
```

```
In [153]: pd.Period('2014', freq='A-DEC') - p
```

```
Out[153]: 7
```





时间区间

- Period_range函数可以构造规则区间序列：

```
In [154]: rng = pd.period_range('2000-01-01', '2000-06-30', freq='M')
```

```
In [155]: rng
```

```
Out[155]: PeriodIndex(['2000-01', '2000-02', '2000-03', '2000-04', '2000-05', '2000-06'], dtype='period[M]', freq='M')
```

```
In [156]: pd.Series(np.random.randn(6), index=rng)
```

```
Out[156]:
```

```
2000-01    -0.514551
```

```
2000-02    -0.559782
```

```
2000-03    -0.783408
```

```
2000-04    -1.797685
```

```
2000-05    -0.172670
```

```
2000-06     0.680215
```

```
Freq: M, dtype: float64
```

Q代表季度为频率,默认的后缀为DEC代表一年以第12个月为结束【最后一个月为12月份】

If you have an array of strings, you can also use the PeriodIndex class:

```
In [157]: values = ['2001Q3', '2002Q2', '2003Q1']
```

```
In [158]: index = pd.PeriodIndex(values, freq='Q-DEC')
```

```
In [159]: index
```

```
Out[159]: PeriodIndex(['2001Q3', '2002Q2', '2003Q1'], dtype='period[Q-DEC]', freq='Q-DEC')
```



区间频率转化

- 使用`asfreq`可以将区间和`PeriodIndex`对象转换为其他的频率。例如，假设我们有一个年度区间，并且想要在一年的开始或结束时将其转换为月度区间：

```
In [160]: p = pd.Period('2007', freq='A-DEC')
```

```
In [161]: p
```

```
Out[161]: Period('2007', 'A-DEC')
```

```
In [162]: p.asfreq('M', how='start')  
Out[162]: Period('2007-01', 'M')
```

```
In [163]: p.asfreq('M', how='end')  
Out[163]: Period('2007-12', 'M')
```

以年为单位，07年6月作为游标。周期的起始点是06年7月，终点为07年6月。

```
In [164]: p = pd.Period('2007', freq='A-JUN')
```

```
In [165]: p
```

```
Out[165]: Period('2007', 'A-JUN')
```

```
In [166]: p.asfreq('M', 'start')  
Out[166]: Period('2006-07', 'M')
```

```
In [167]: p.asfreq('M', 'end')  
Out[167]: Period('2007-06', 'M')
```



区间频率转化

```
In [170]: rng = pd.period_range('2006', '2009', freq='A-DEC')
```

```
In [171]: ts = pd.Series(np.random.randn(len(rng)), index=rng)
```

```
In [172]: ts
```

```
Out[172]:
```

```
2006    1.607578
```

```
2007    0.200381
```

```
2008   -0.834068
```

```
2009   -0.302988
```

```
Freq: A-DEC, dtype: float64
```

```
In [173]: ts.asfreq('M', how='start')
```

```
Out[173]:
```

```
2006-01    1.607578
```

```
2007-01    0.200381
```

```
2008-01   -0.834068
```

```
2009-01   -0.302988
```

```
Freq: M, dtype: float64
```

```
In [174]: ts.asfreq('B', how='end')
```

```
Out[174]:
```

```
2006-12-29    1.607578
```

```
2007-12-31    0.200381
```

```
2008-12-31   -0.834068
```

```
2009-12-31   -0.302988
```

```
Freq: B, dtype: float64
```





时间戳和区间之间的转化

```
In [188]: rng = pd.date_range('2000-01-01', periods=3, freq='M')
```

```
In [189]: ts = pd.Series(np.random.randn(3), index=rng)
```

```
In [190]: ts
```

```
Out[190]:
```

```
2000-01-31    1.663261
```

```
2000-02-29   -0.996206
```

```
2000-03-31    1.521760
```

```
Freq: M, dtype: float64
```

```
In [191]: pts = ts.to_period()
```

```
In [192]: pts
```

```
Out[192]:
```

```
2000-01    1.663261
```

```
2000-02   -0.996206
```

```
2000-03    1.521760
```

```
Freq: M, dtype: float64
```

```
In [193]: rng = pd.date_range('1/29/2000', periods=6, freq='D')
```

```
In [194]: ts2 = pd.Series(np.random.randn(6), index=rng)
```

```
In [195]: ts2
```

```
Out[195]:
```

```
2000-01-29    0.244175
```

```
2000-01-30    0.423331
```

```
2000-01-31   -0.654040
```

```
2000-02-01    2.089154
```

```
2000-02-02   -0.060220
```

```
2000-02-03   -0.167933
```

```
Freq: D, dtype: float64
```

```
In [196]: ts2.to_period('M')
```

```
Out[196]:
```

```
2000-01    0.244175
```

```
2000-01    0.423331
```

```
2000-01   -0.654040
```

```
2000-02    2.089154
```

```
2000-02   -0.060220
```

```
2000-02   -0.167933
```

```
Freq: M, dtype: float64
```





时间戳和区间之间的转化

To convert back to timestamps, use `to_timestamp`:

```
In [197]: pts = ts2.to_period()
```

```
In [198]: pts
```

```
Out[198]:
```

```
2000-01-29    0.244175  
2000-01-30    0.423331  
2000-01-31    -0.654040  
2000-02-01    2.089154  
2000-02-02    -0.060220  
2000-02-03    -0.167933
```

```
Freq: D, dtype: float64
```

```
In [199]: pts.to_timestamp(how='end')
```

```
Out[199]:
```

```
2000-01-29    0.244175  
2000-01-30    0.423331  
2000-01-31    -0.654040  
2000-02-01    2.089154  
2000-02-02    -0.060220  
2000-02-03    -0.167933
```

```
Freq: D, dtype: float64
```

```
In [195]: ts2  
Out[195]:  
2000-01-29    0.244175  
2000-01-30    0.423331  
2000-01-31    -0.654040  
2000-02-01    2.089154  
2000-02-02    -0.060220  
2000-02-03    -0.167933  
Freq: D, dtype: float64
```





时间序列的重采样 (resampling)





重采样 (Resampling)

- 重采样: 将时间序列从一个频率转换为另一个频率的过程。
 - 高频率数据聚合到低频率: 下采样 (Downsampling)
 - 低频率数据转换到高频率: 上采样 (Upsampling)
 - 其他: 如每周三的数据转化为每周五的数据



重采样 (Resampling)

```
In [208]: rng = pd.date_range('2000-01-01', periods=100, freq='D')
```

```
In [209]: ts = pd.Series(np.random.randn(len(rng)), index=rng)
```

```
In [210]: ts
```

```
Out[210]:
```

```
2000-01-01    0.631634  
2000-01-02   -1.594313  
2000-01-03   -1.519937  
2000-01-04    1.108752  
2000-01-05    1.255853  
2000-01-06   -0.024330  
2000-01-07   -2.047939  
2000-01-08   -0.272657  
2000-01-09   -1.692615  
2000-01-10    1.423830  
...  
2000-03-31   -0.007852
```

```
2000-04-01   -1.638806  
2000-04-02    1.401227  
2000-04-03    1.758539  
2000-04-04    0.628932  
2000-04-05   -0.423776  
2000-04-06    0.789740  
2000-04-07    0.937568  
2000-04-08   -2.253294  
2000-04-09   -1.772919
```

```
Freq: D, Length: 100, dtype: float64
```

```
In [211]: ts.resample('M').mean()
```

```
Out[211]:
```

```
2000-01-31   -0.165893  
2000-02-29    0.078606  
2000-03-31    0.223811  
2000-04-30   -0.063643
```

```
Freq: M, dtype: float64
```

```
In [212]: ts.resample('M', kind='period').mean()
```

```
Out[212]:
```

```
2000-01   -0.165893  
2000-02    0.078606  
2000-03    0.223811  
2000-04   -0.063643
```

```
Freq: M, dtype: float64
```





下采样 (Downsampling)

默认情况下，左箱体边界是包含的，即00:00的值包含在了00:00-00:05内；但传递close='right'后，间隔的闭合端改为了右边。

```
In [213]: rng = pd.date_range('2000-01-01', periods=12, freq='T')
```

```
In [214]: ts = pd.Series(np.arange(12), index=rng)
```

```
In [215]: ts
```

```
Out[215]:
```

```
2000-01-01 00:00:00    0  
2000-01-01 00:01:00    1  
2000-01-01 00:02:00    2  
2000-01-01 00:03:00    3  
2000-01-01 00:04:00    4  
2000-01-01 00:05:00    5  
2000-01-01 00:06:00    6  
2000-01-01 00:07:00    7  
2000-01-01 00:08:00    8  
2000-01-01 00:09:00    9  
2000-01-01 00:10:00   10  
2000-01-01 00:11:00   11  
  
Freq: T, dtype: int64
```

```
In [216]: ts.resample('5min', closed='right').sum()
```

```
Out[216]:
```

```
1999-12-31 23:55:00    0  
2000-01-01 00:00:00    15  
2000-01-01 00:05:00    40  
2000-01-01 00:10:00    11  
  
Freq: 5T, dtype: int64
```

```
In [218]: ts.resample('5min', closed='right', label='right').sum()
```

```
Out[218]:
```

```
2000-01-01 00:00:00    0  
2000-01-01 00:05:00    15  
2000-01-01 00:10:00    40  
2000-01-01 00:15:00    11  
  
Freq: 5T, dtype: int64
```

默认情况下，产生的时
间序列按照每个箱体左
边的时间戳标记；传递
label='right'后，将使用右
箱体边界时间序列。





下采样 (Downsampling)

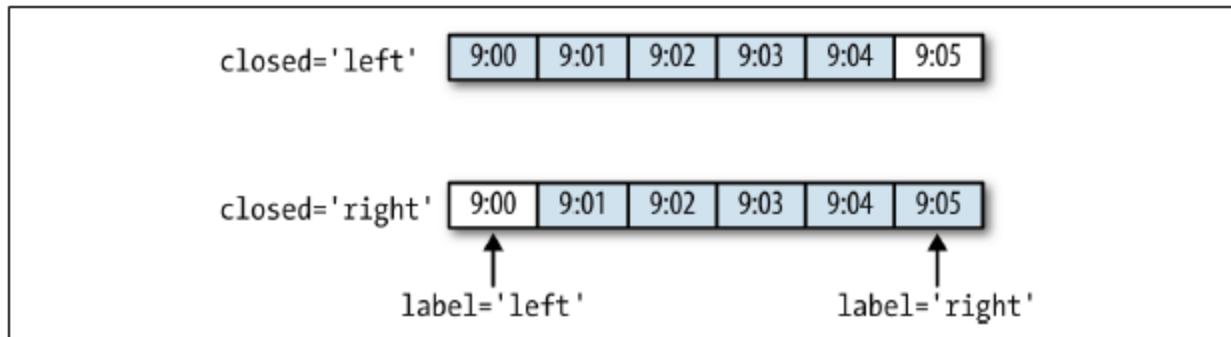


Figure 11-3. Five-minute resampling illustration of closed, label conventions

```
In [219]: ts.resample('5min', closed='right',
.....:                      label='right', loffset='-1s').sum()
```

Out[219]:

```
1999-12-31 23:59:59      0
2000-01-01 00:04:59     15
2000-01-01 00:09:59     40
2000-01-01 00:14:59     11
Freq: 5T, dtype: int64
```





Resample的方法参数

Argument	Description
freq	String or DateOffset indicating desired resampled frequency (e.g., 'M', '5min', or Second(15))
axis	Axis to resample on; default axis=0
fill_method	How to interpolate when upsampling, as in 'ffill' or 'bfill'; by default does no interpolation
closed	In downsampling, which end of each interval is closed (inclusive), 'right' or 'left'
label	In downsampling, how to label the aggregated result, with the 'right' or 'left' bin edge (e.g., the 9:30 to 9:35 five-minute interval could be labeled 9:30 or 9:35)
loffset	Time adjustment to the bin labels, such as '-1s' / Second(-1) to shift the aggregate labels one second earlier
limit	When forward or backward filling, the maximum number of periods to fill
kind	Aggregate to periods ('period') or timestamps ('timestamp'); defaults to the type of index the time series has
convention	When resampling periods, the convention ('start' or 'end') for converting the low-frequency period to high frequency; defaults to 'end'





【扩展】开端-峰值-谷值-结束 (OHLC) 重采样

- 金融领域：为每个数据箱计算四个值是一种流行的时间序列聚合方法。使用OHLC聚合函数。

In [215]: ts

Out[215]:

```
2000-01-01 00:00:00    0
2000-01-01 00:01:00    1
2000-01-01 00:02:00    2
2000-01-01 00:03:00    3
2000-01-01 00:04:00    4
2000-01-01 00:05:00    5
2000-01-01 00:06:00    6
2000-01-01 00:07:00    7
2000-01-01 00:08:00    8
2000-01-01 00:09:00    9
2000-01-01 00:10:00   10
2000-01-01 00:11:00   11
```

Freq: T, dtype: int64

Open-High-Low-Close (OHLC) resampling

1

```
ts.resample('5min').ohlc()
```

		open	high	low	close
	2000-01-01 00:00:00	0	4	0	4
	2000-01-01 00:05:00	5	9	5	9
	2000-01-01 00:10:00	10	11	10	11





上采样 (Upsampling) 和差值 (Interpolation)

```
In [221]: frame = pd.DataFrame(np.random.randn(2, 4),
.....                         index=pd.date_range('1/1/2000', periods=2,
.....                               freq='W-WED'),
.....                         columns=['Colorado', 'Texas', 'New York', 'Ohio'])
```

```
In [222]: frame
Out[222]:
          Colorado      Texas   New York      Ohio
2000-01-05 -0.896431  0.677263  0.036503  0.087102
2000-01-12 -0.046662  0.927238  0.482284 -0.867130
```

```
In [223]: df_daily = frame.resample('D').asfreq()
```

```
In [224]: df_daily
Out[224]:
          Colorado      Texas   New York      Ohio
2000-01-05 -0.896431  0.677263  0.036503  0.087102
2000-01-06      NaN       NaN       NaN       NaN
2000-01-07      NaN       NaN       NaN       NaN
2000-01-08      NaN       NaN       NaN       NaN
2000-01-09      NaN       NaN       NaN       NaN
2000-01-10      NaN       NaN       NaN       NaN
2000-01-11      NaN       NaN       NaN       NaN
2000-01-12 -0.046662  0.927238  0.482284 -0.867130
```





In [222]: frame

Out[222]:

		Colorado	Texas	New York	Ohio
2000-01-05	-0.896431	0.677263	0.036503	0.087102	
2000-01-12	-0.046662	0.927238	0.482284	-0.867130	

In [225]: frame.resample('D').ffill()

Out[225]:

	Colorado	Texas	New York	Ohio
2000-01-05	-0.896431	0.677263	0.036503	0.087102
2000-01-06	-0.896431	0.677263	0.036503	0.087102
2000-01-07	-0.896431	0.677263	0.036503	0.087102
2000-01-08	-0.896431	0.677263	0.036503	0.087102
2000-01-09	-0.896431	0.677263	0.036503	0.087102
2000-01-10	-0.896431	0.677263	0.036503	0.087102
2000-01-11	-0.896431	0.677263	0.036503	0.087102
2000-01-12	-0.046662	0.927238	0.482284	-0.867130

In [226]: frame.resample('D').ffill(limit=2)

Out[226]:

	Colorado	Texas	New York	Ohio
2000-01-05	-0.896431	0.677263	0.036503	0.087102
2000-01-06	-0.896431	0.677263	0.036503	0.087102
2000-01-07	-0.896431	0.677263	0.036503	0.087102
2000-01-08		NaN	NaN	NaN
2000-01-09		NaN	NaN	NaN
2000-01-10		NaN	NaN	NaN
2000-01-11		NaN	NaN	NaN
2000-01-12	-0.046662	0.927238	0.482284	-0.867130



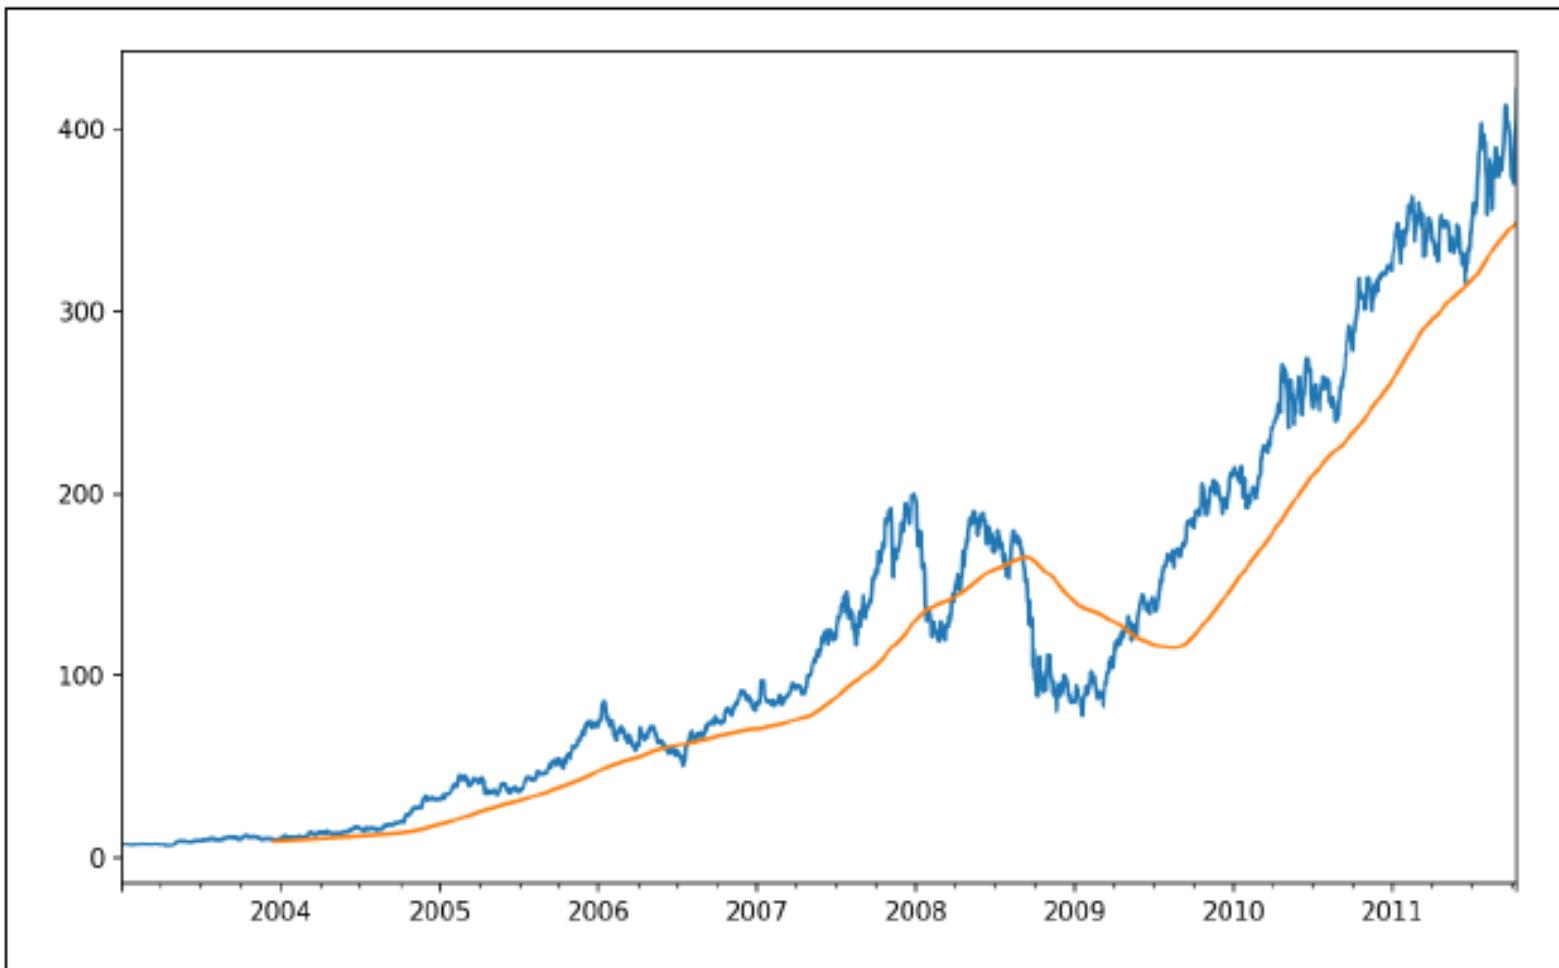
移动窗口函数 (Moving window functions)





移动窗口函数

可用于平滑噪声或粗糙的数据！





移动窗口函数

```
In [235]: close_px_all = pd.read_csv('examples/stock_px_2.csv',
.....:                                     parse_dates=True, index_col=0)
```

```
In [236]: close_px = close_px_all[['AAPL', 'MSFT', 'XOM']]
```

```
In [237]: close_px = close_px.resample('B').ffill()
```

1 close_px

	AAPL	MSFT	XOM
2003-01-02	7.40	21.11	29.22
2003-01-03	7.45	21.14	29.24
2003-01-06	7.45	21.52	29.96
2003-01-07	7.43	21.93	28.95
2003-01-08	7.28	21.31	28.83
...
2011-10-10	388.81	26.94	76.28
2011-10-11	400.29	27.00	76.27
2011-10-12	402.19	26.96	77.16
2011-10-13	408.43	27.18	76.37
2011-10-14	422.00	27.27	78.11

2292 rows × 3 columns



移动窗口函数

```
In [238]: close_px.AAPL.plot()
```

```
Out[238]: <matplotlib.axes._subplots.AxesSubplot at 0x7f2f2570cf98>
```

```
In [239]: close_px.AAPL.rolling(250).mean().plot()
```

根据250日滑动窗口分组





移动窗口函数

```
In [241]: appl_std250 = close_px.AAPL.rolling(250, min_periods=10).std()
```

```
In [242]: appl_std250[5:12]
```

```
Out[242]:
```

```
2003-01-09      NaN  
2003-01-10      NaN  
2003-01-13      NaN  
2003-01-14      NaN  
2003-01-15    0.077496  
2003-01-16    0.074760  
2003-01-17    0.112368
```

```
Freq: B, Name: AAPL, dtype: float64
```

```
In [243]: appl_std250.plot()
```

时间序列的起始位置少于窗口区间，所以有NaN

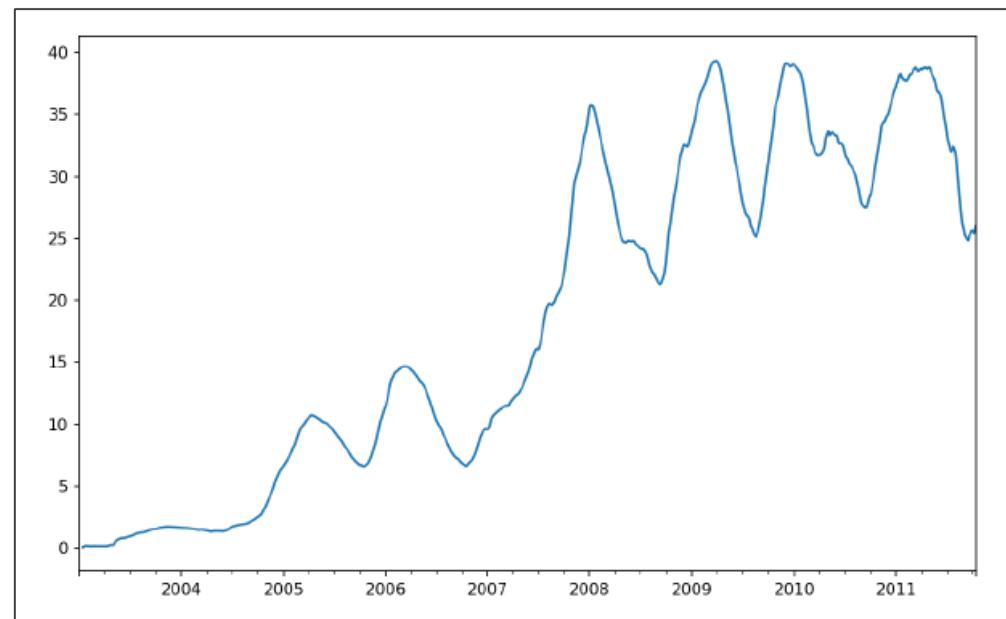


Figure 11-5. Apple 250-day daily return standard deviation



移动窗口函数

- 为了扩展窗口均值，可以使用expanding算子。扩展均值从时间序列起始位置开始时间窗口，并增加窗口的大小，直到它涵盖整个序列。

```
In [244]: expanding_mean = appl_std250.expanding().mean()
```

```
In [246]: close_px.rolling(60).mean().plot(logy=True)
```

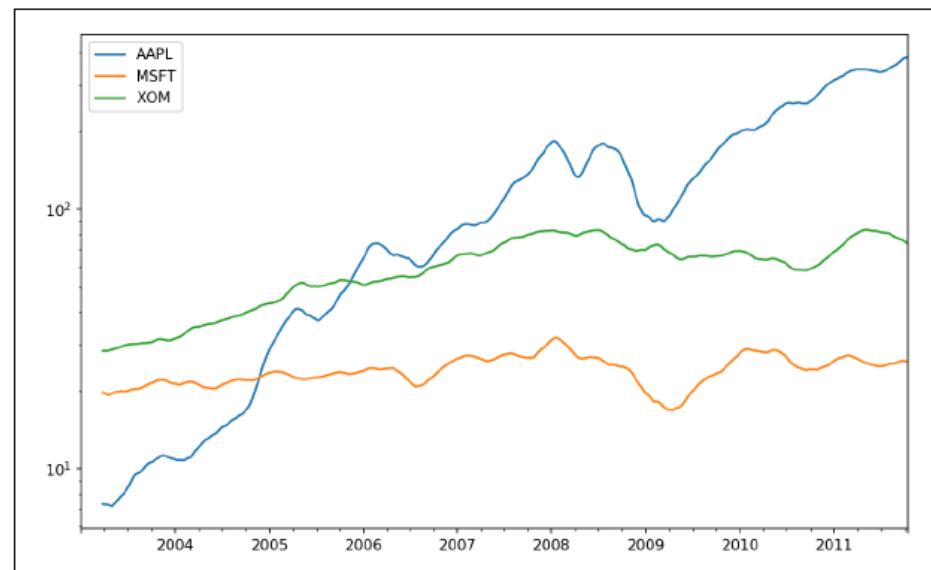


Figure 11-6. Stocks prices 60-day MA (log Y-axis)



其他滑动窗口函数

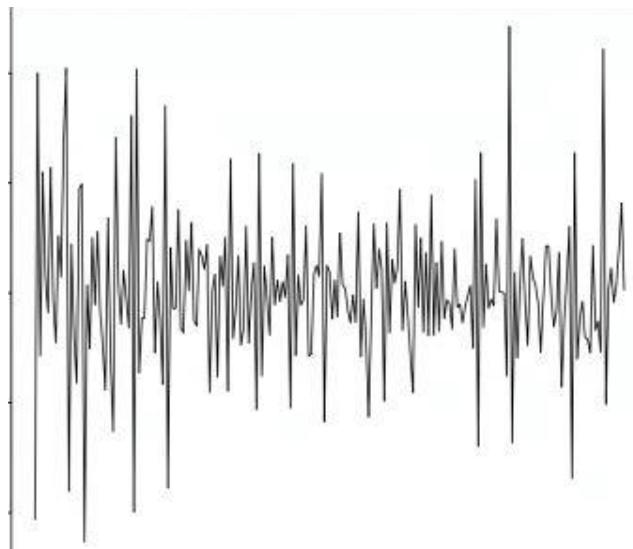
- 指数型滑动平均函数
- 二元移动窗口函数
- 用户自定义的移动窗口函数（使用apply方法等）
-



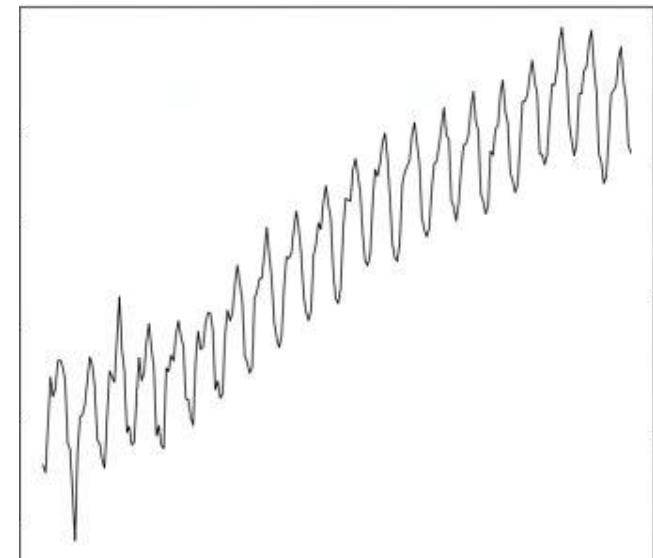


时间序列分析的后续内容

- 时间序列的预处理
- 平稳时间序列的分析（自相关、差分运算）
- 非平稳时间序列的分析
- 预测
- ...



看起来平稳

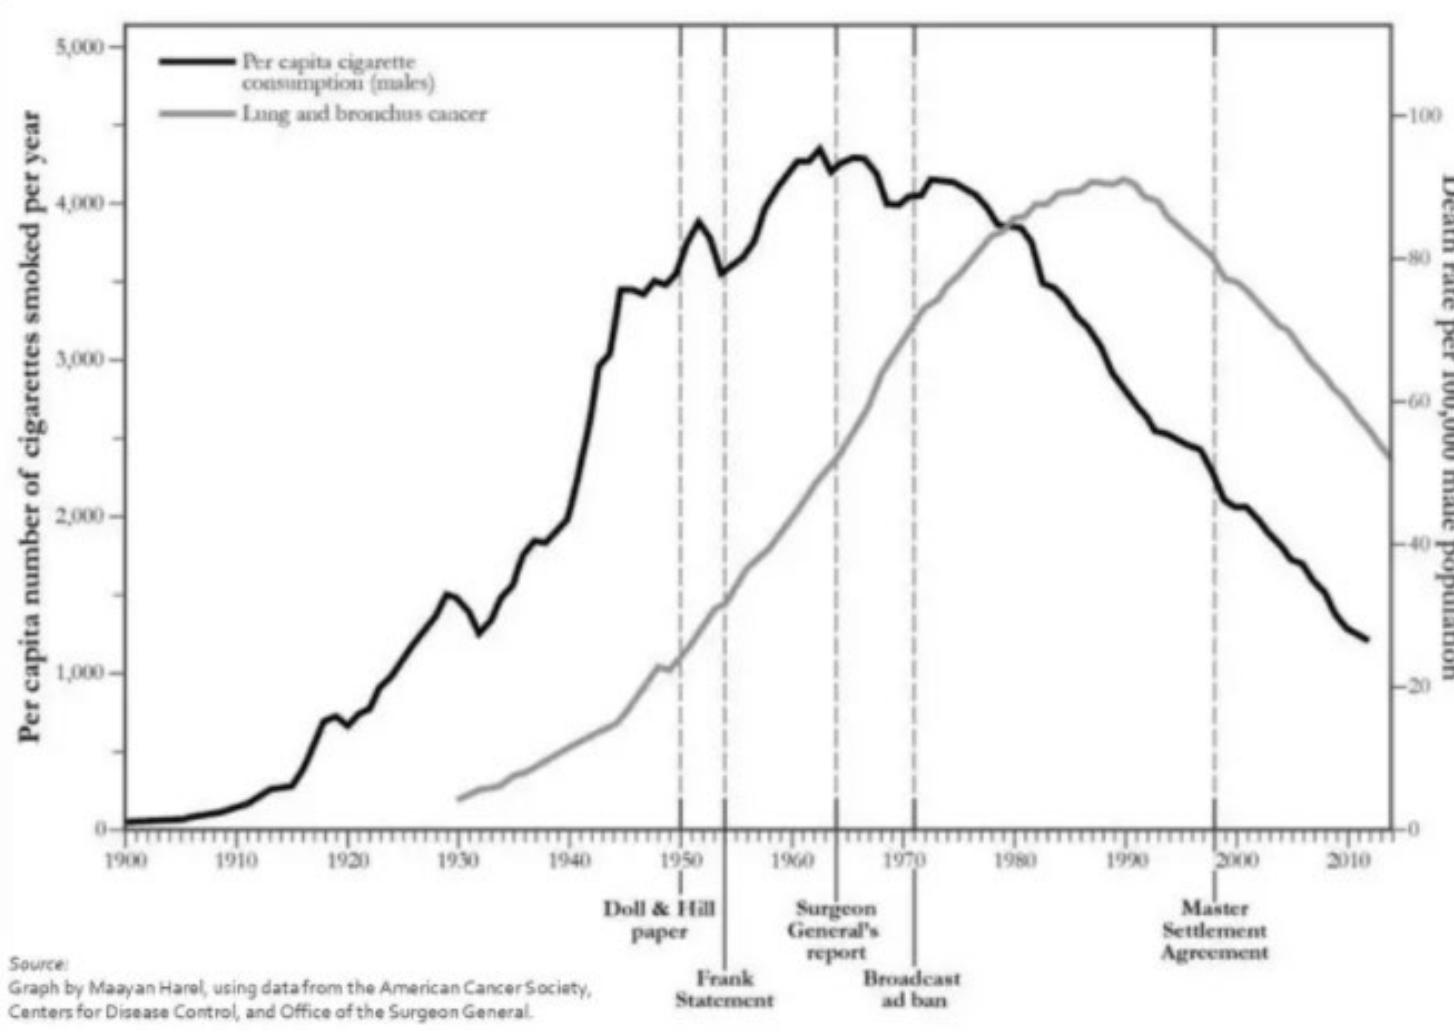


看起来不平稳





时间序列分析的后续内容





个人作业4（共15分）

- 第一部分：时间序列分析
- 第二部分：图像数据分析
- 第三部分：文本数据分析
- 预计5月31日布置





课堂练习





谢谢！

